Deep Learning

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

Overview

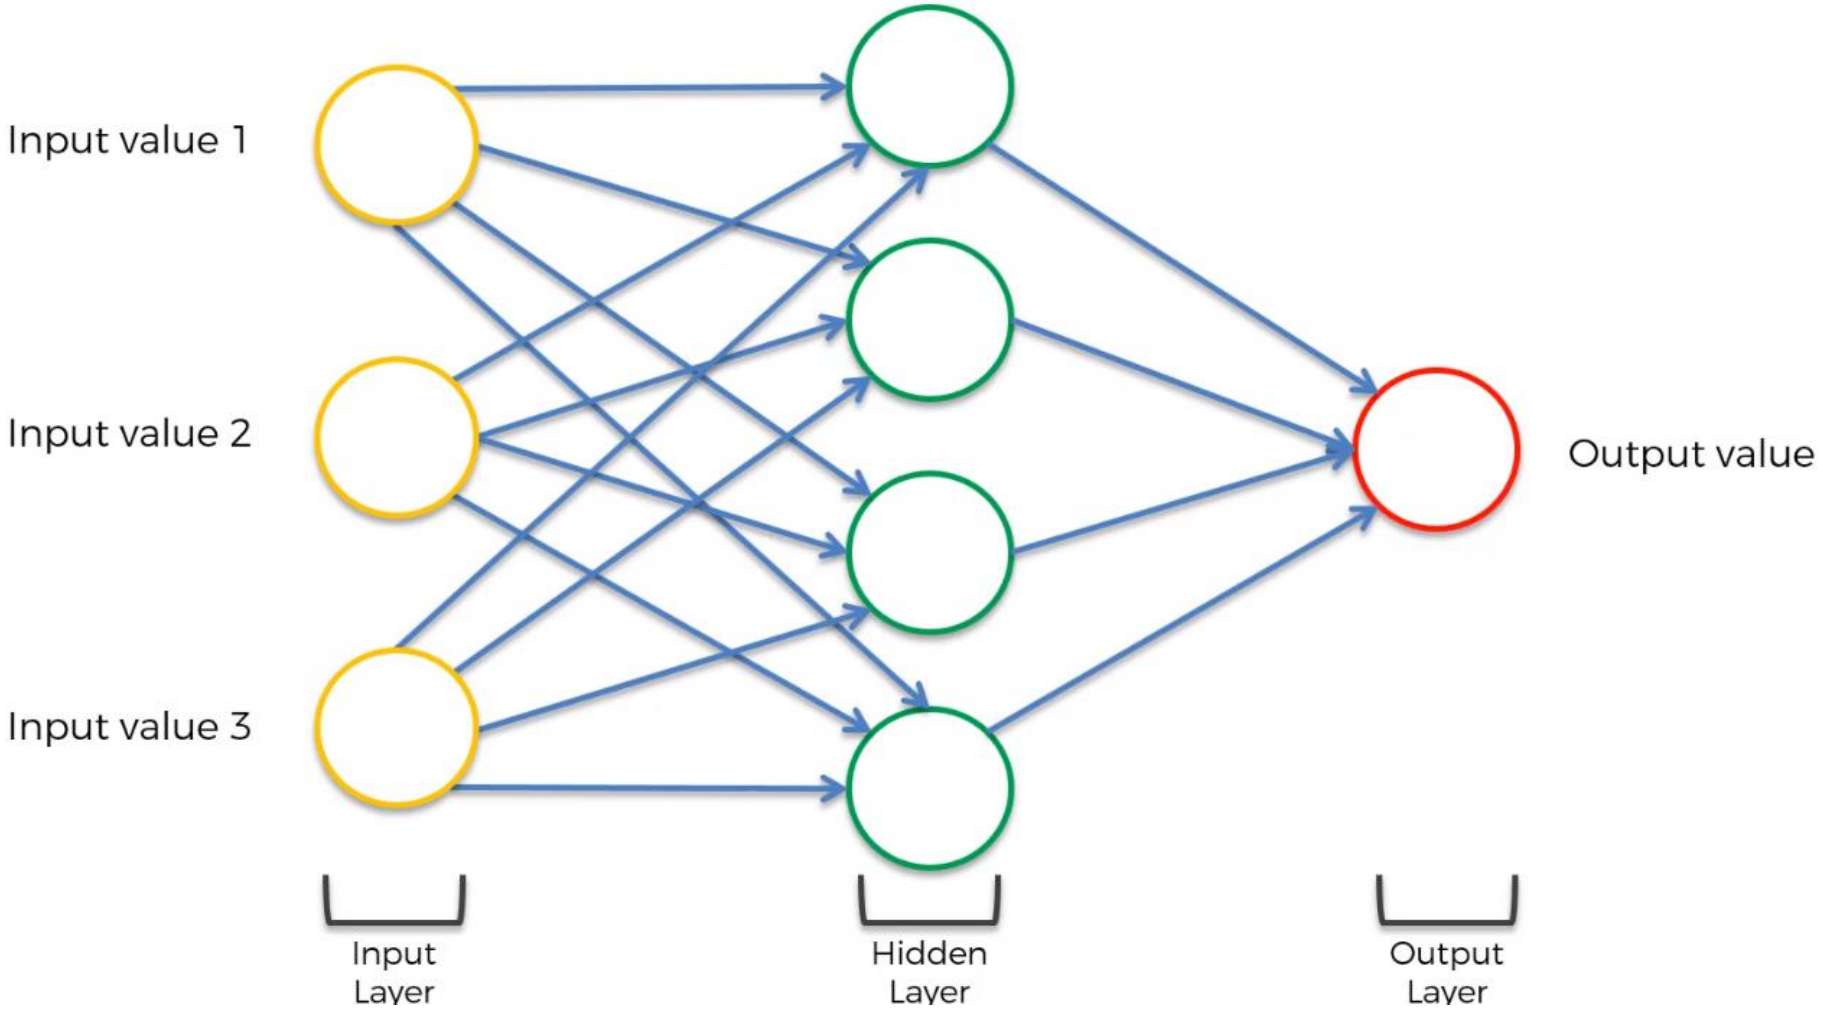

https://www.udemy.com/course/machinelearning/

Overview

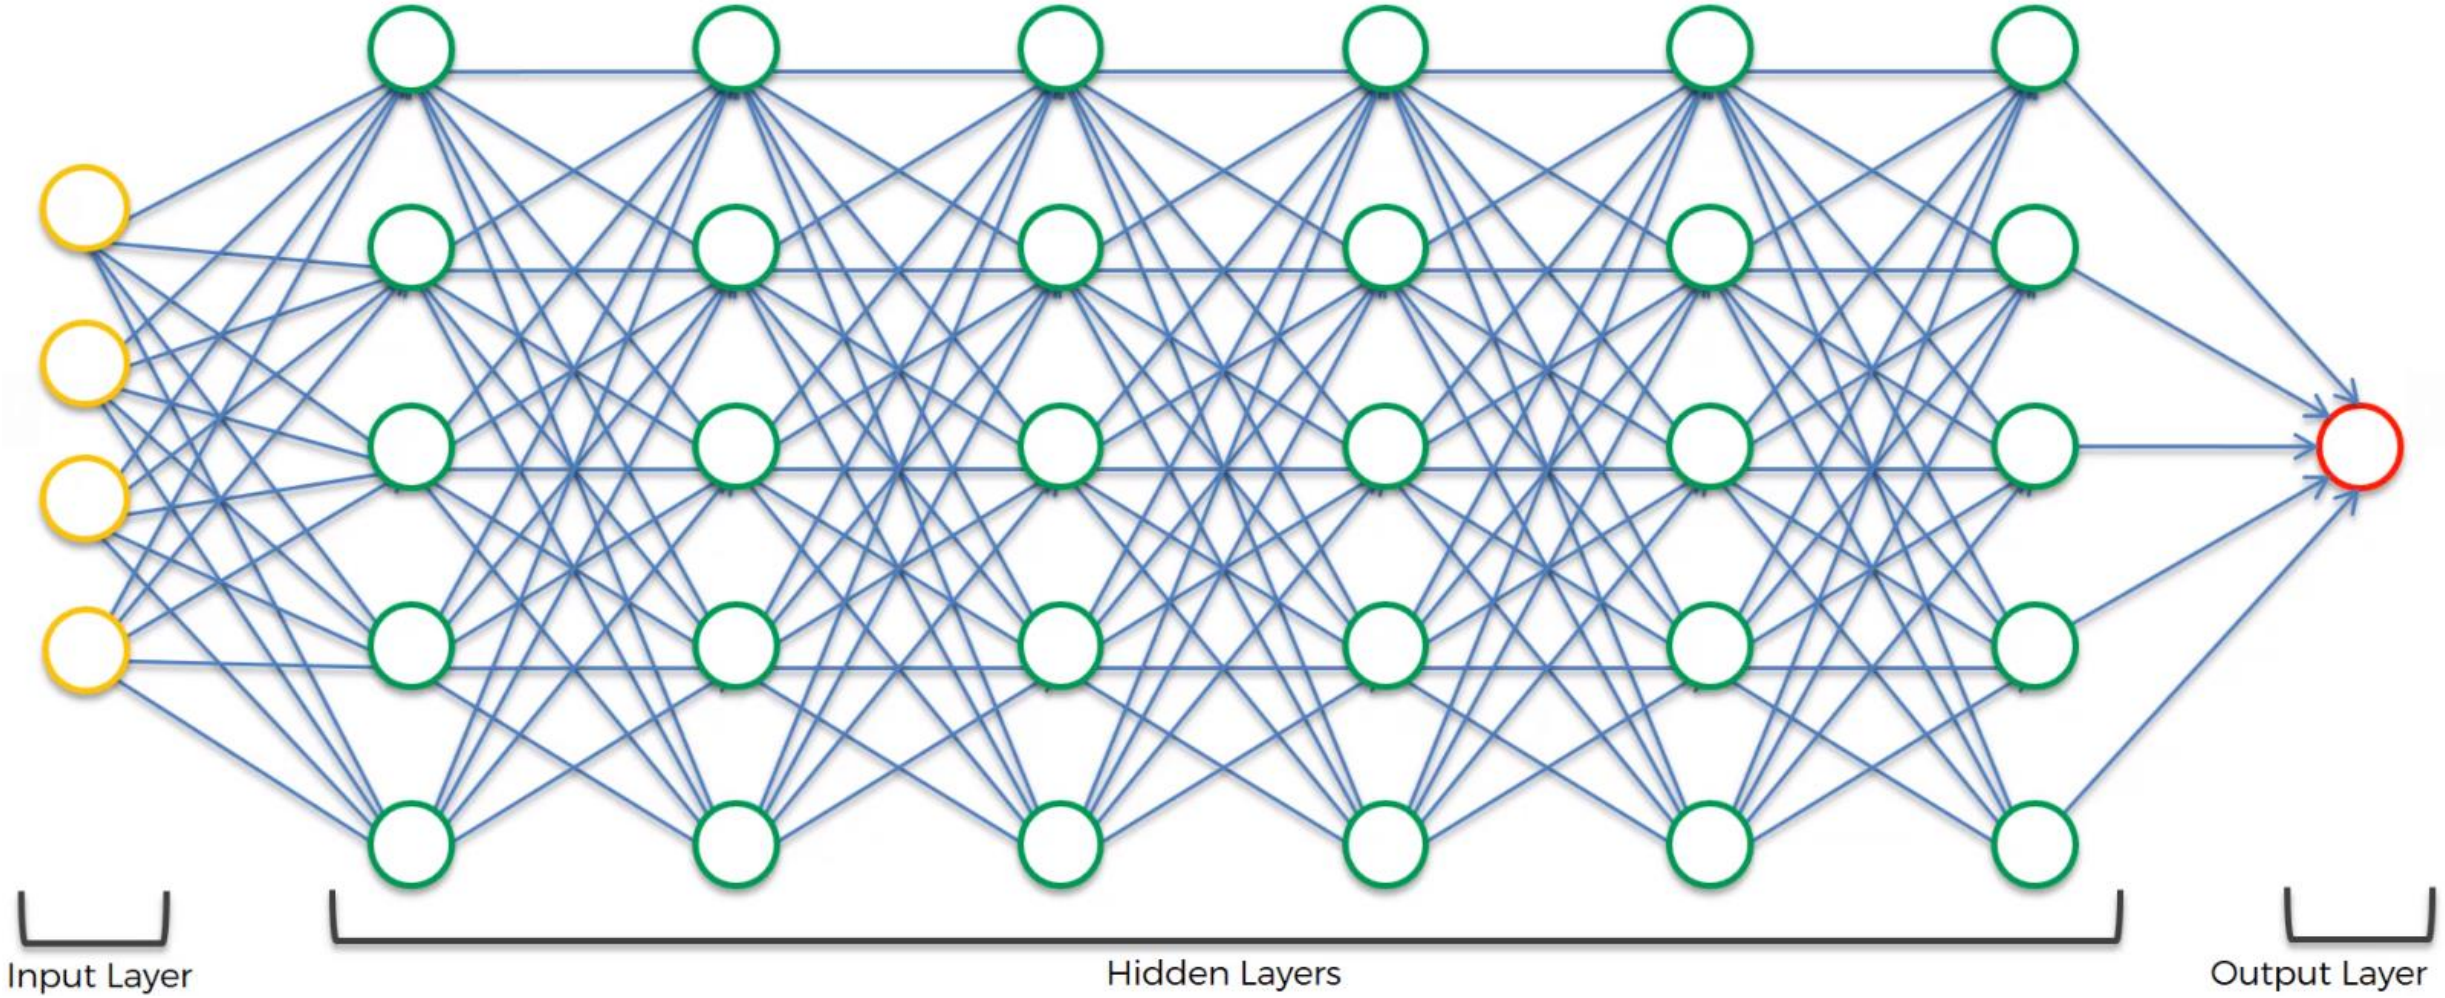

https://www.udemy.com/course/machinelearning/

Overview

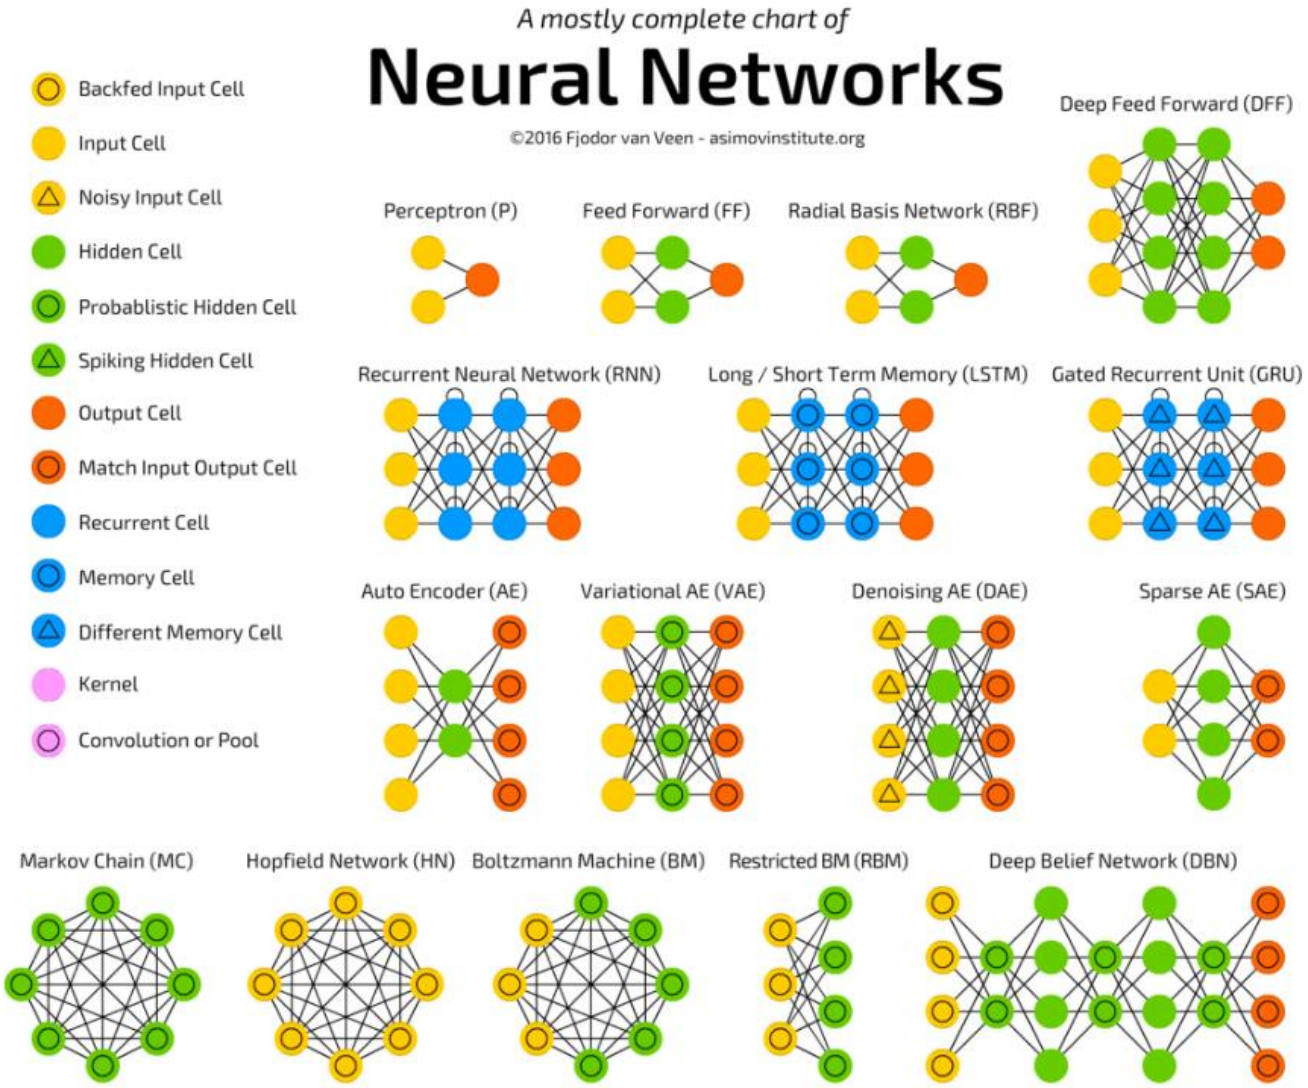

https://towardsdatascience.com/the-mostly-complete-chart-of-neural-networks-explained-3fb6f2367464

Overview

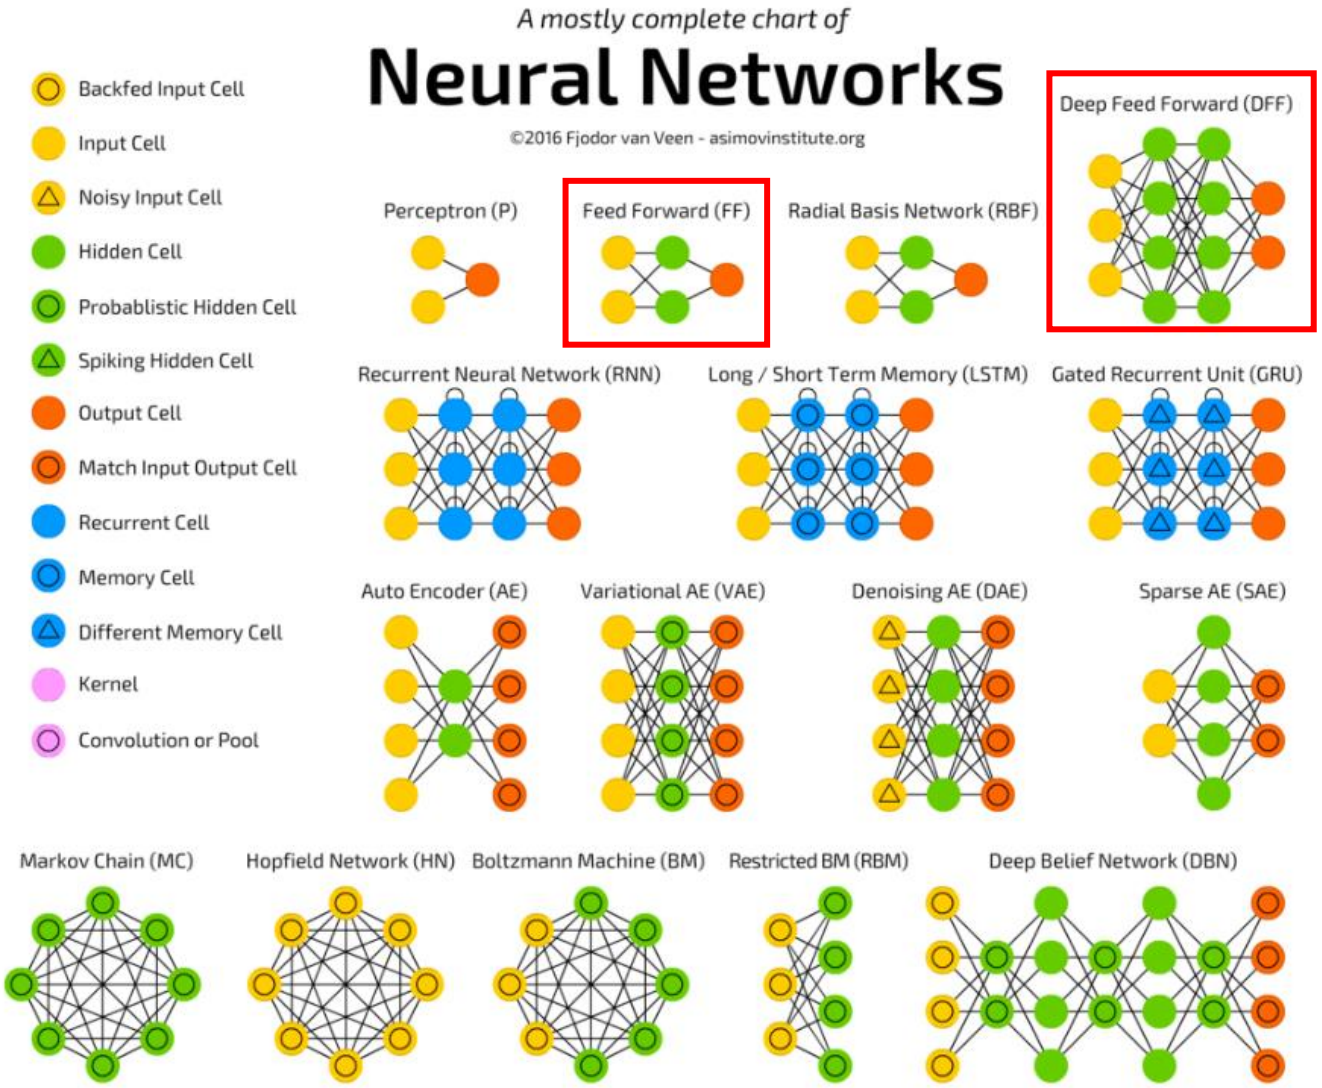

https://towardsdatascience.com/the-mostly-complete-chart-of-neural-networks-explained-3fb6f2367464

Overview

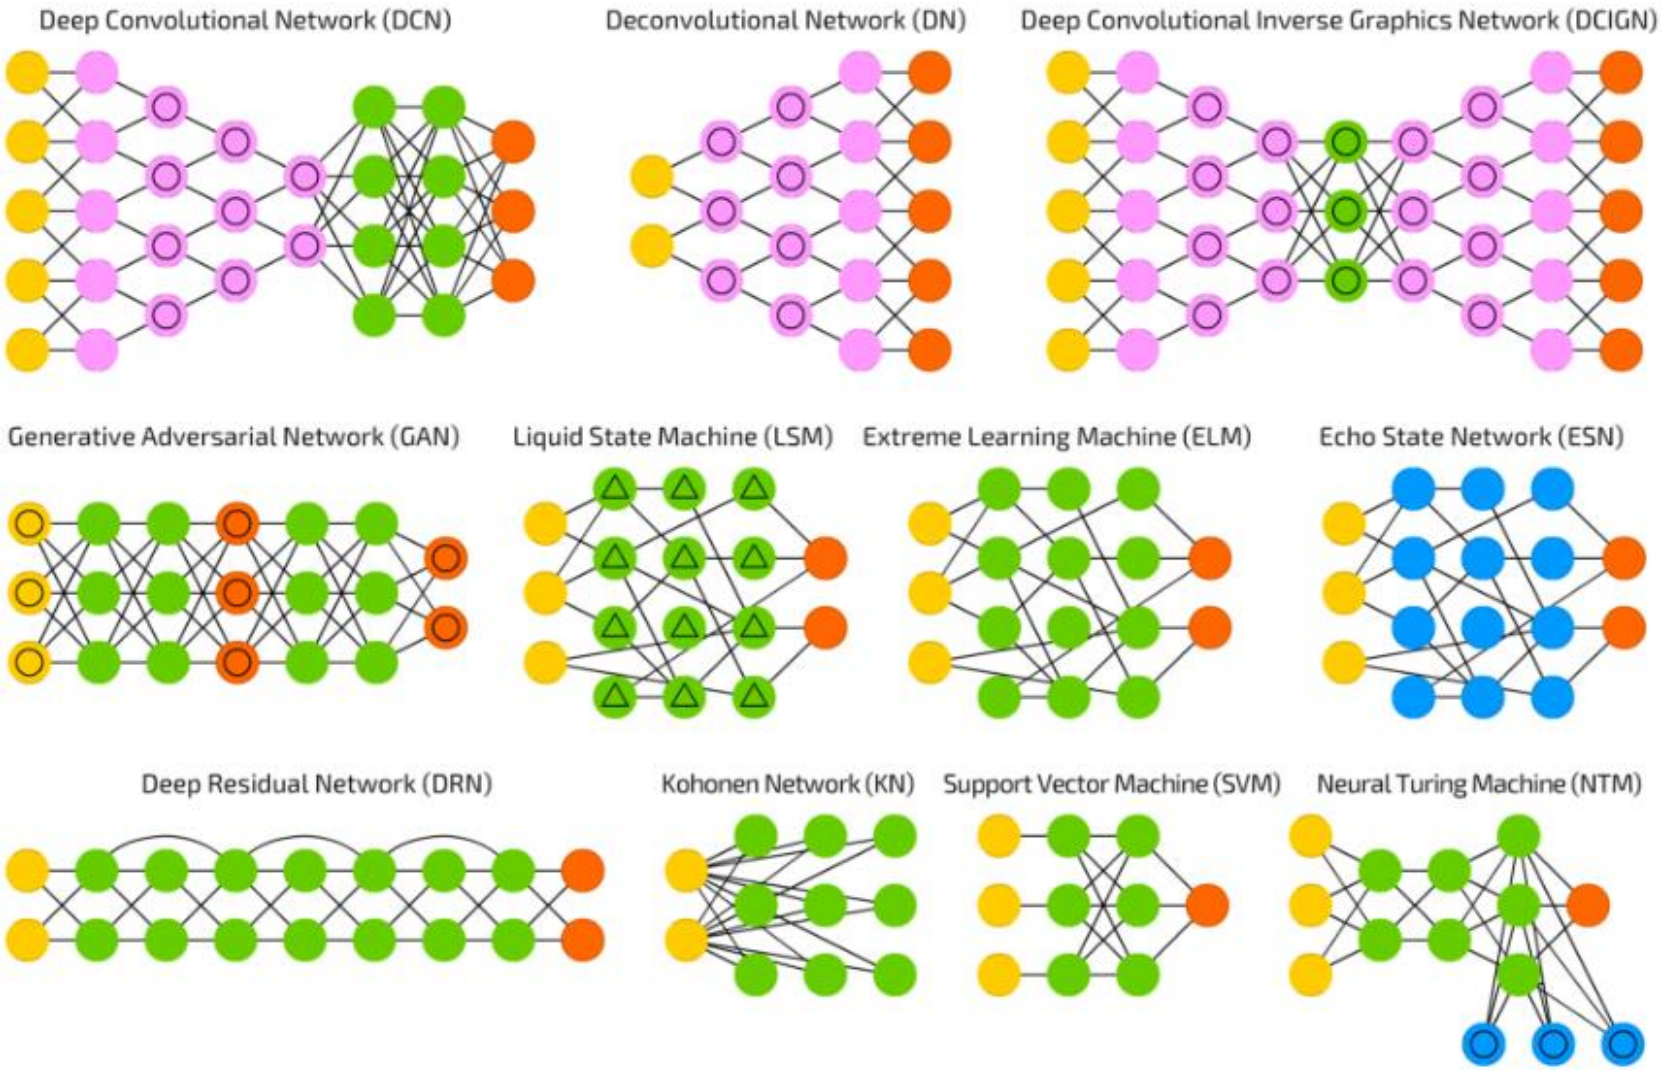

https://towardsdatascience.com/the-mostly-complete-chart-of-neural-networks-explained-3fb6f2367464

Overview

- **Supervised Learning**: CNNs classifying images in imagenet.
- **Unsupervised Learning**: Boltzmann Machines, AutoEncoders, GANs, DC-GANS, VAE, SOMs, etc.
- **Reinforcement**: Deep Convolutional Q-Learning that plays videogames from pixel input, AlphaGO, etc.

Overview

- Artificial Neural Networks for Regression and Classification
- Convolutional Neural Networks for Computer Vision
- Recurrent Neural Networks for Time Series Analysis
- Self Organizing Maps for Feature Extraction
- Deep Boltzmann Machines for Recommendation Systems
- Auto Encoders for Recommendation Systems

https://www.udemy.com/course/machinelearning/

Overview

- **Artificial Neural Networks for Regression and Classification**
- **Convolutional Neural Networks for Computer Vision**
- Recurrent Neural Networks for Time Series Analysis
- Self Organizing Maps for Feature Extraction
- Deep Boltzmann Machines for Recommendation Systems
- Auto Encoders for Recommendation Systems

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation

https://www.udemy.com/course/machinelearning/

The Neuron

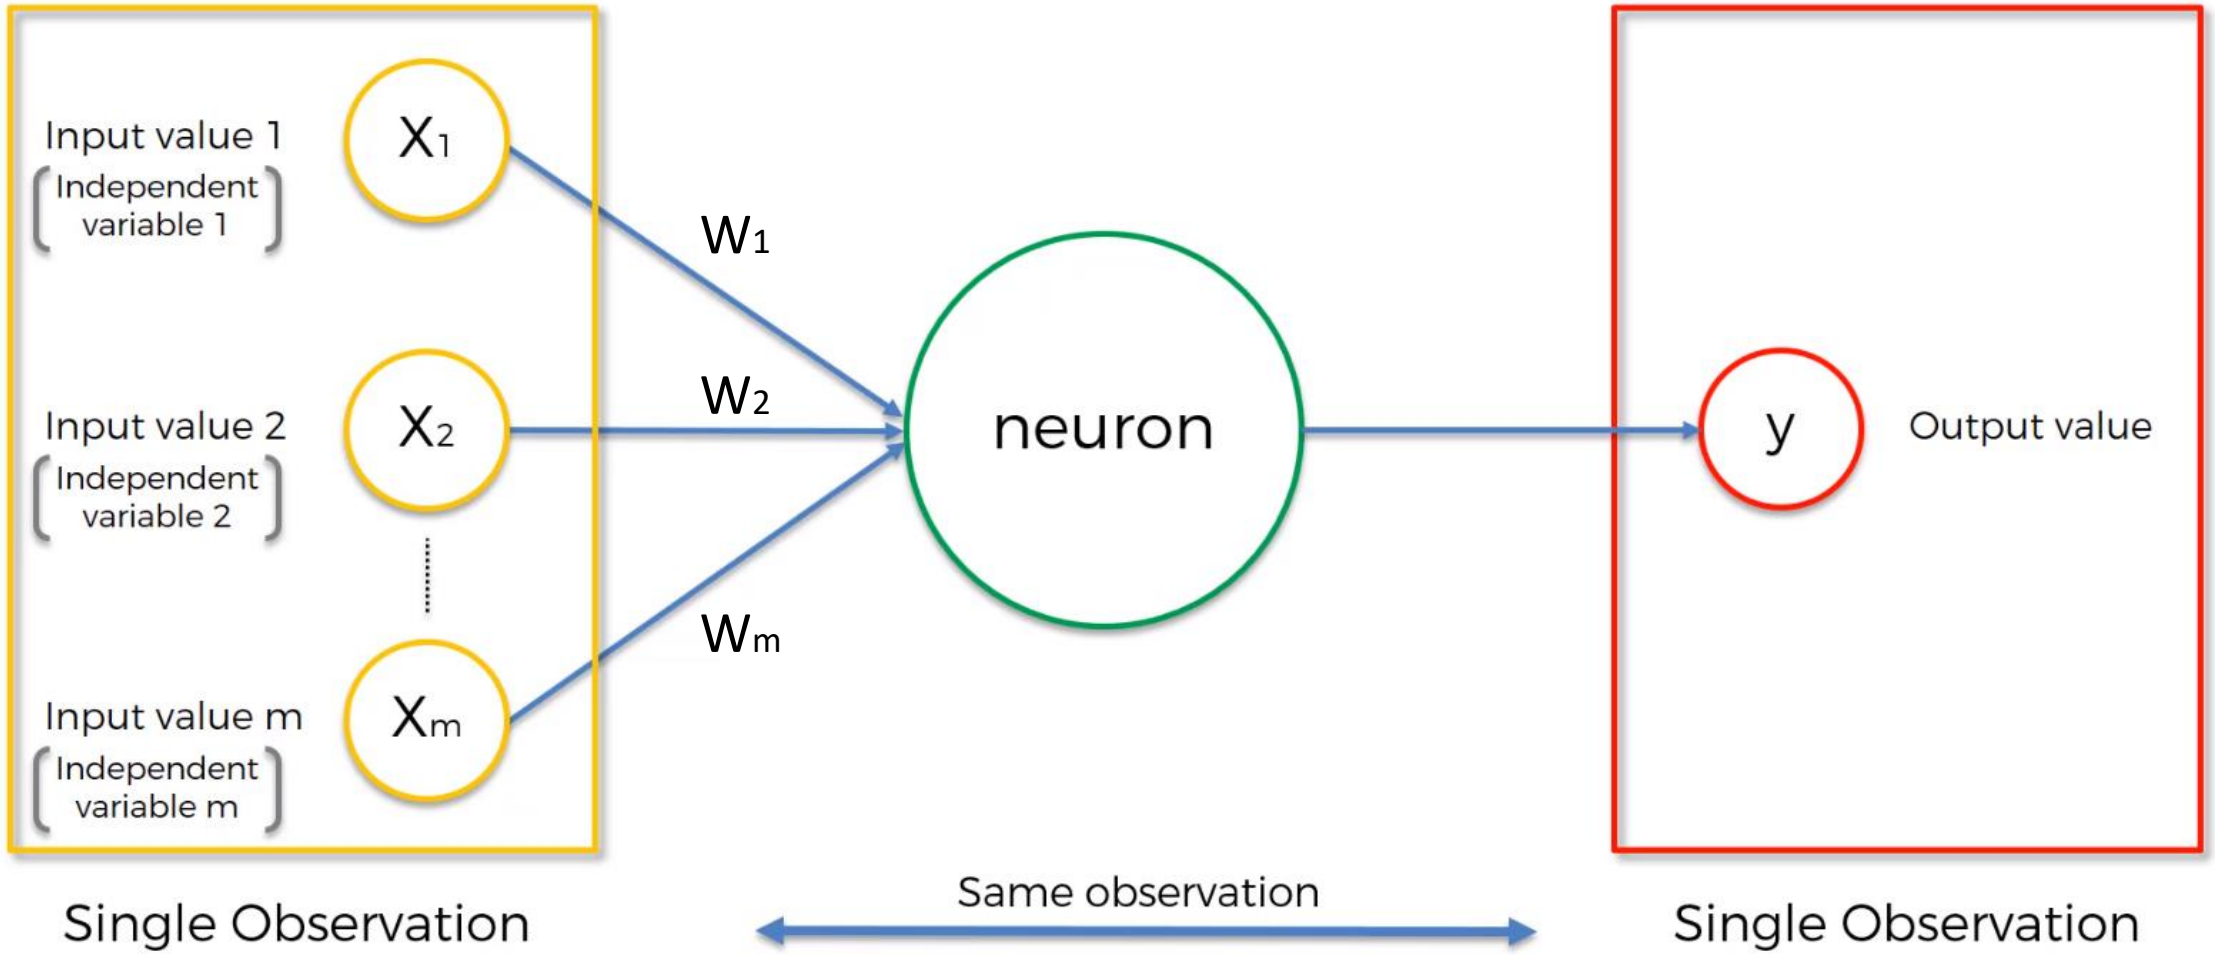

https://www.udemy.com/course/machinelearning/

The Neuron

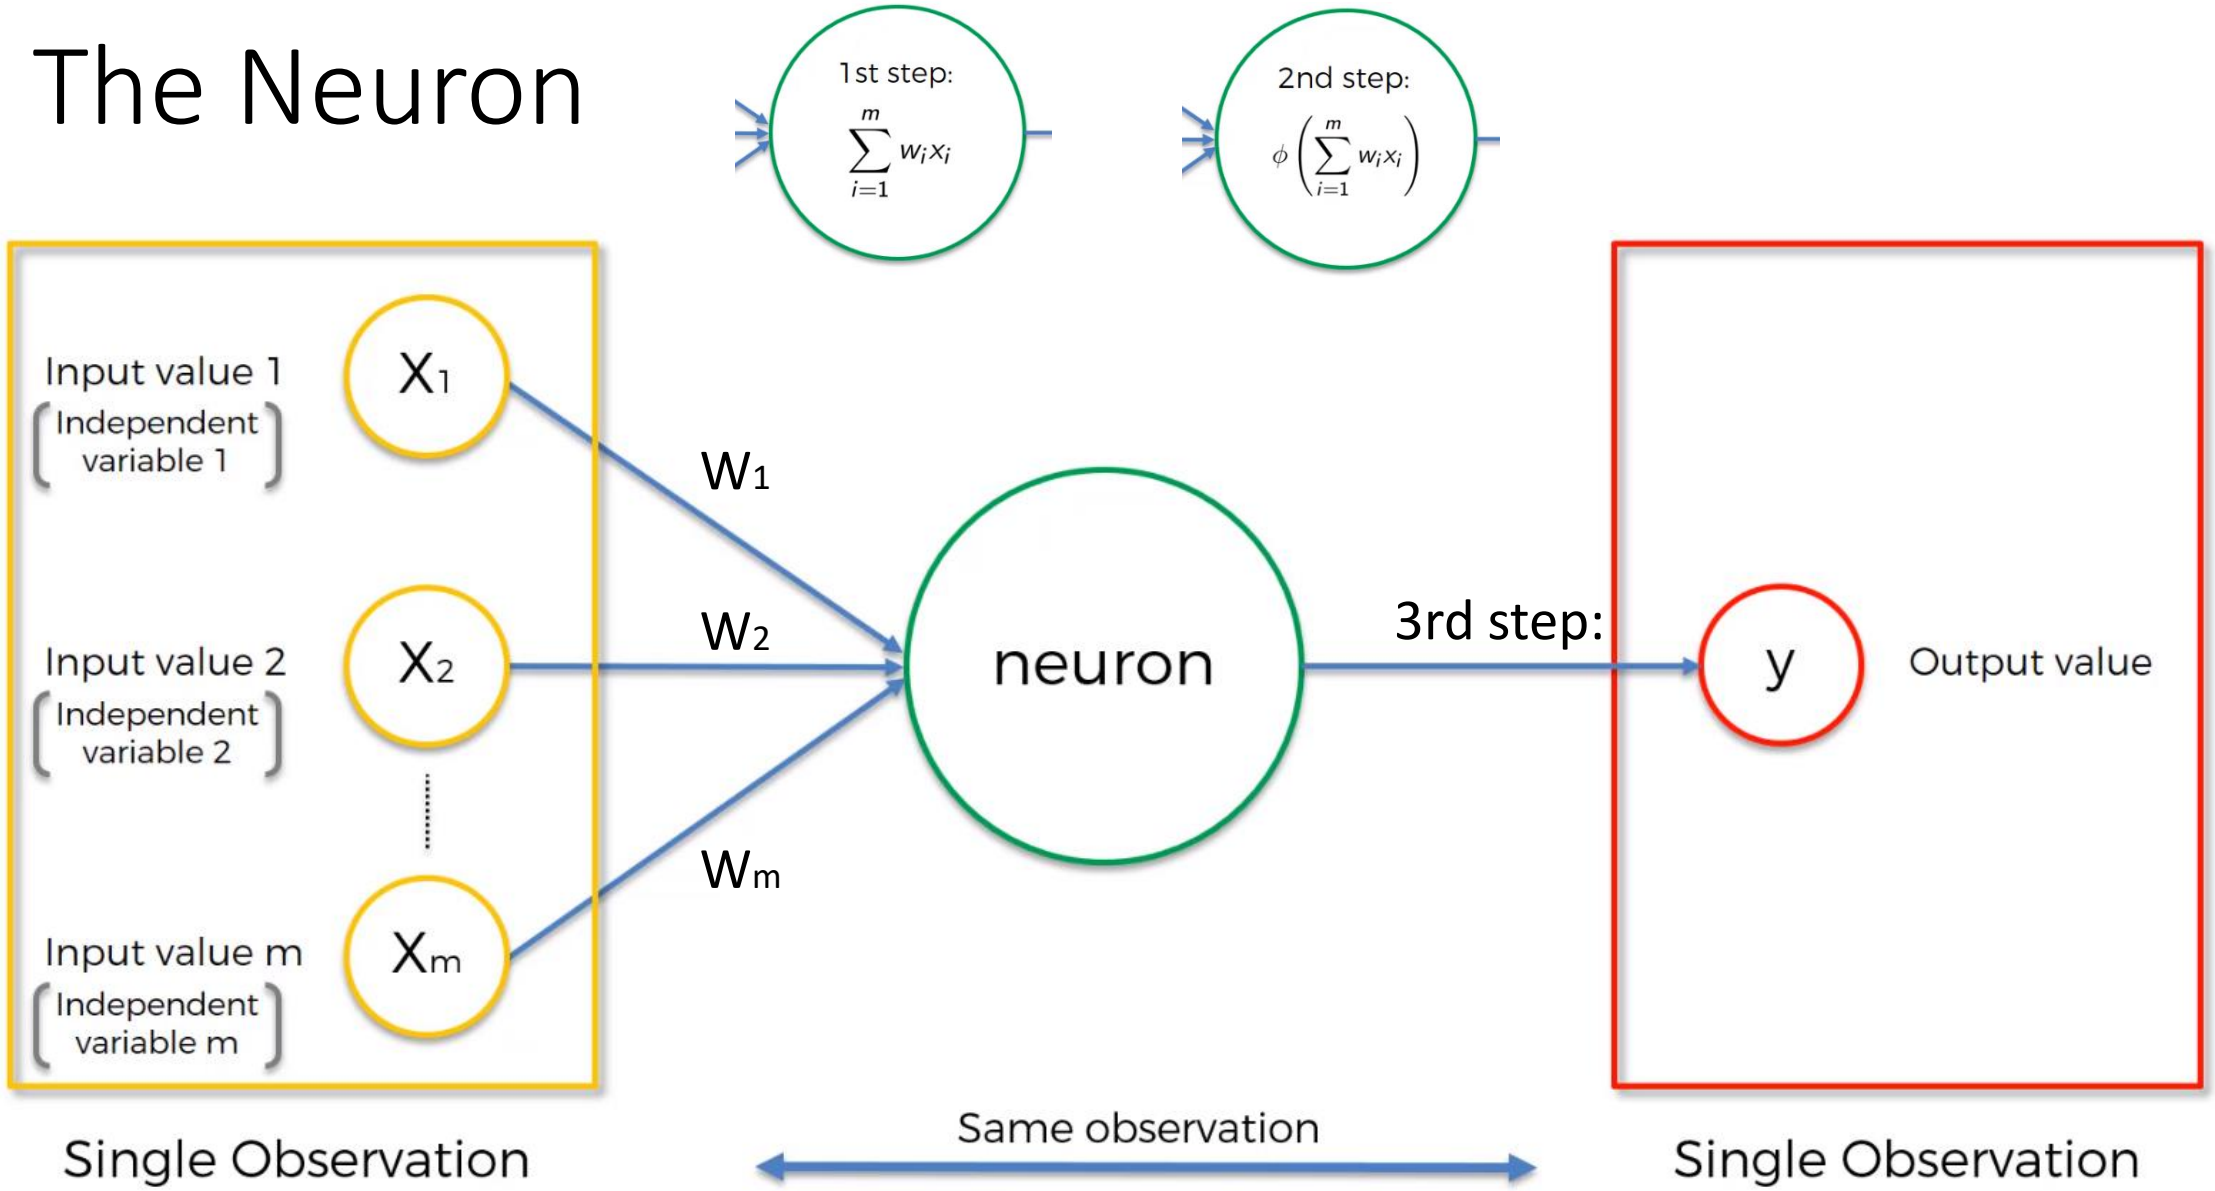

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

The Activation Function

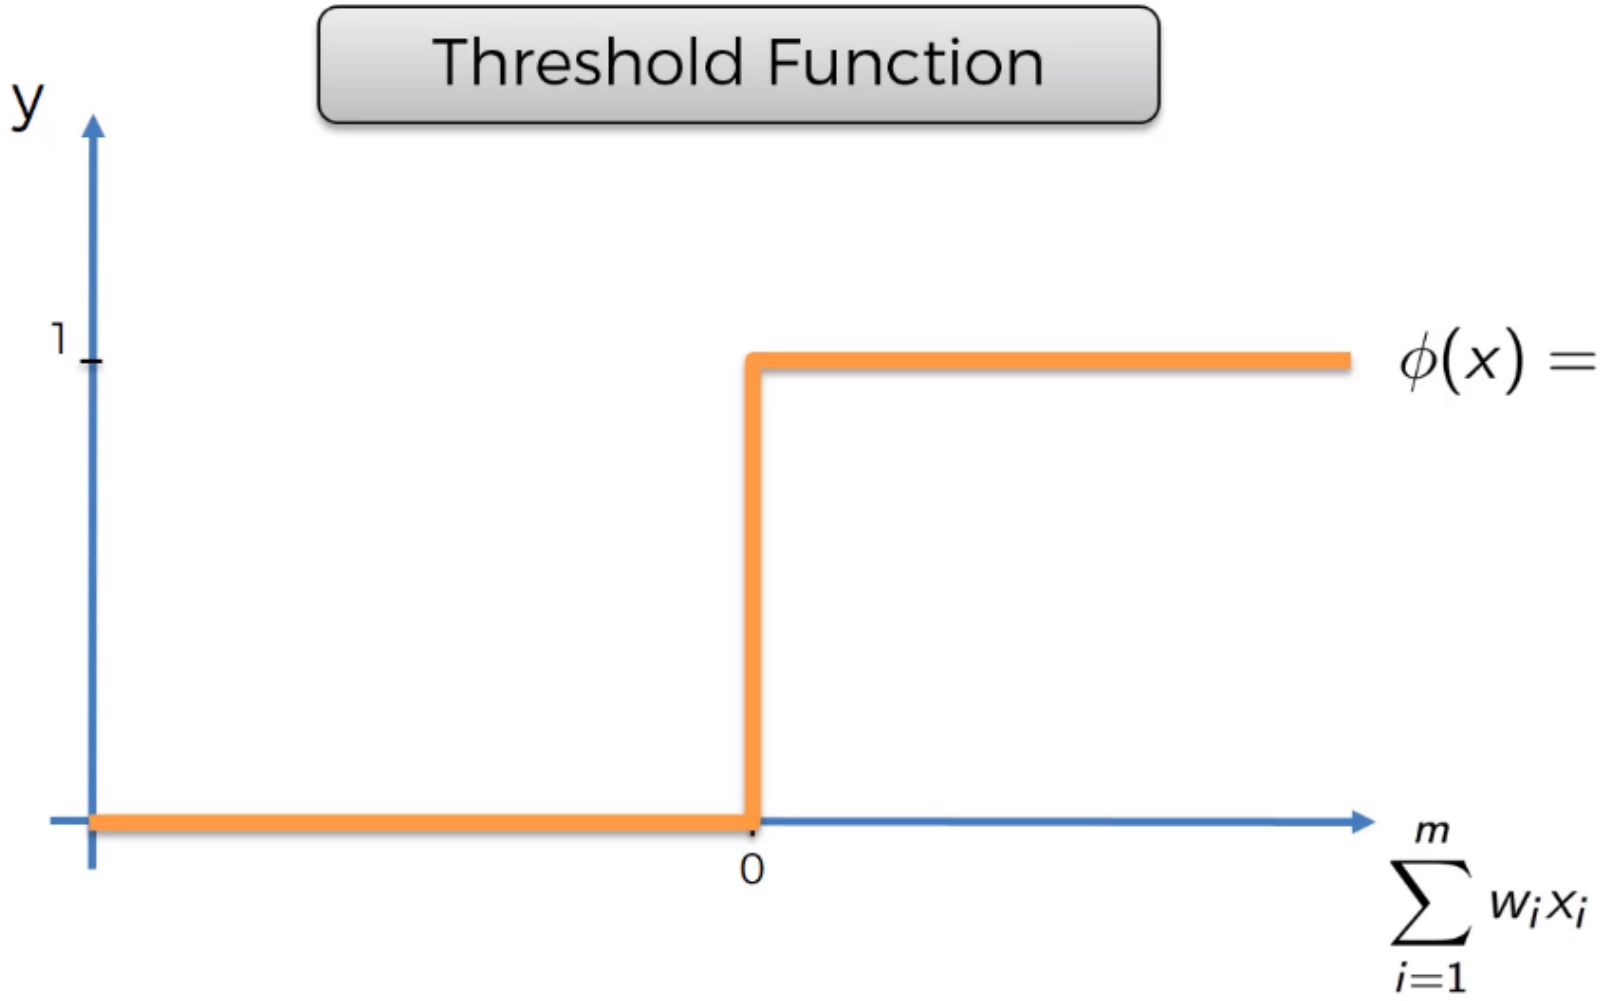

$$\phi(x) = \begin{cases} 1 \text{ if } x \geq 0 \\ 0 \text{ if } x < 0 \end{cases}$$

https://www.udemy.com/course/machinelearning/

The Activation Function

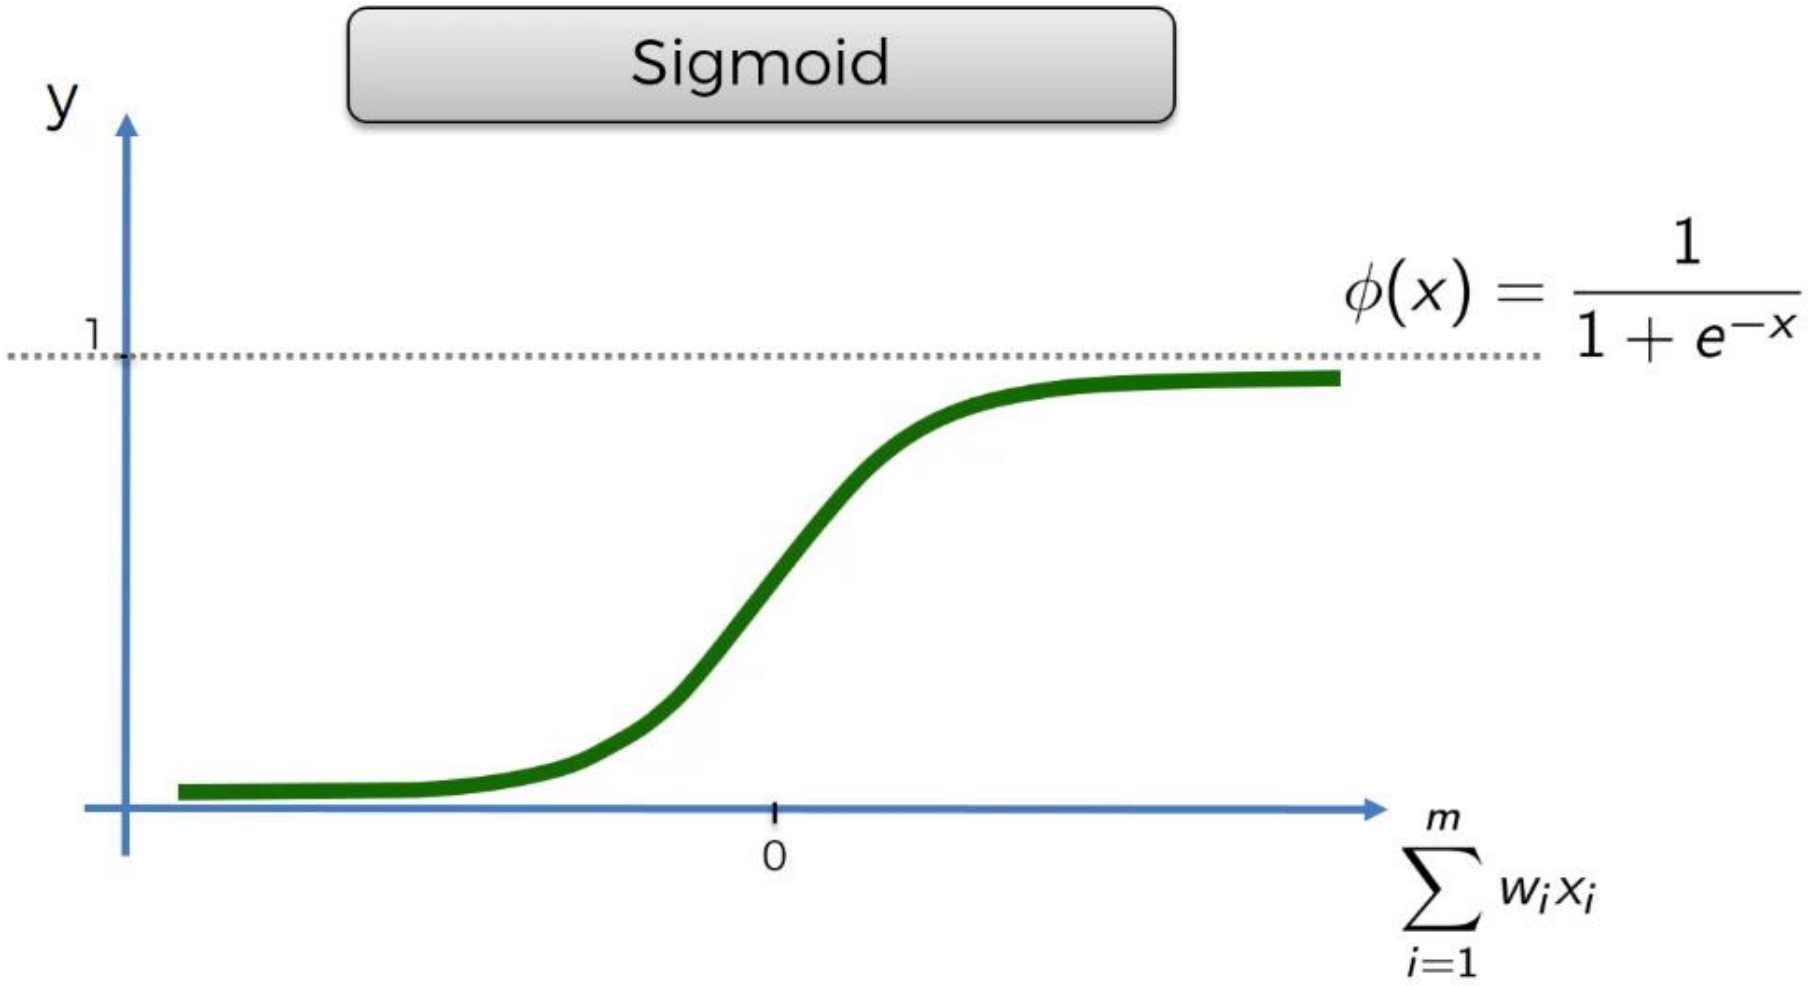

https://www.udemy.com/course/machinelearning/

The Activation Function

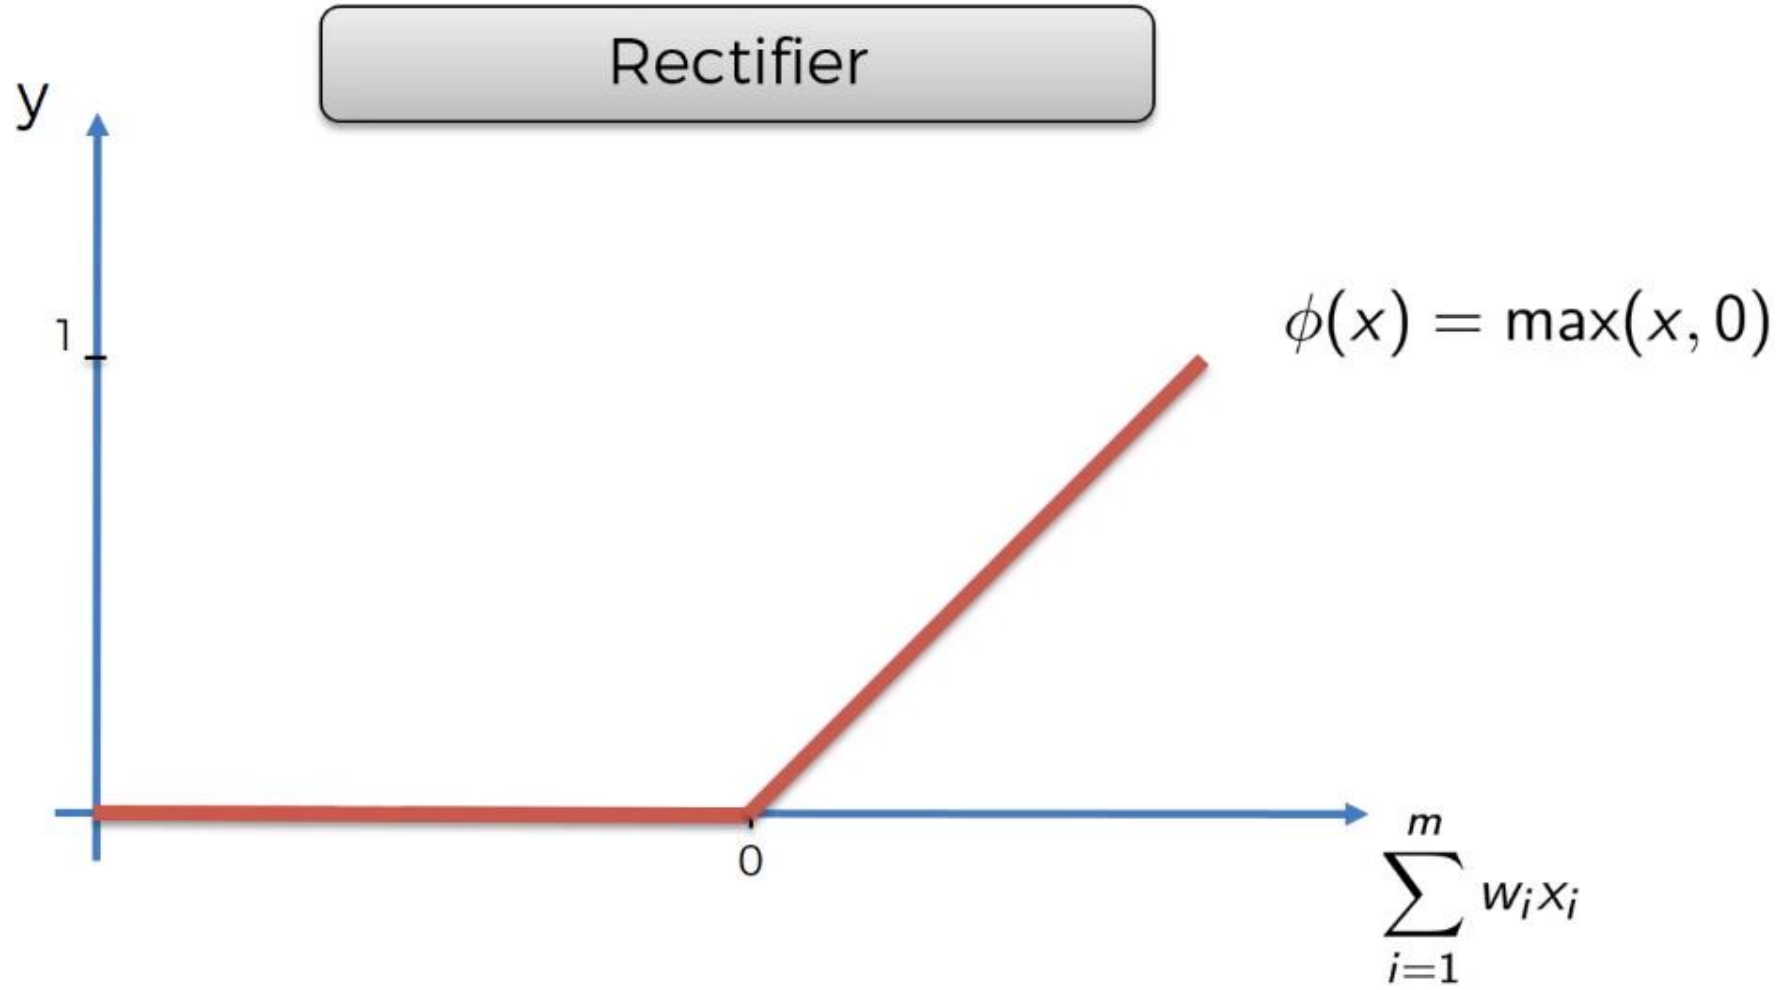

https://www.udemy.com/course/machinelearning/

The Activation Function

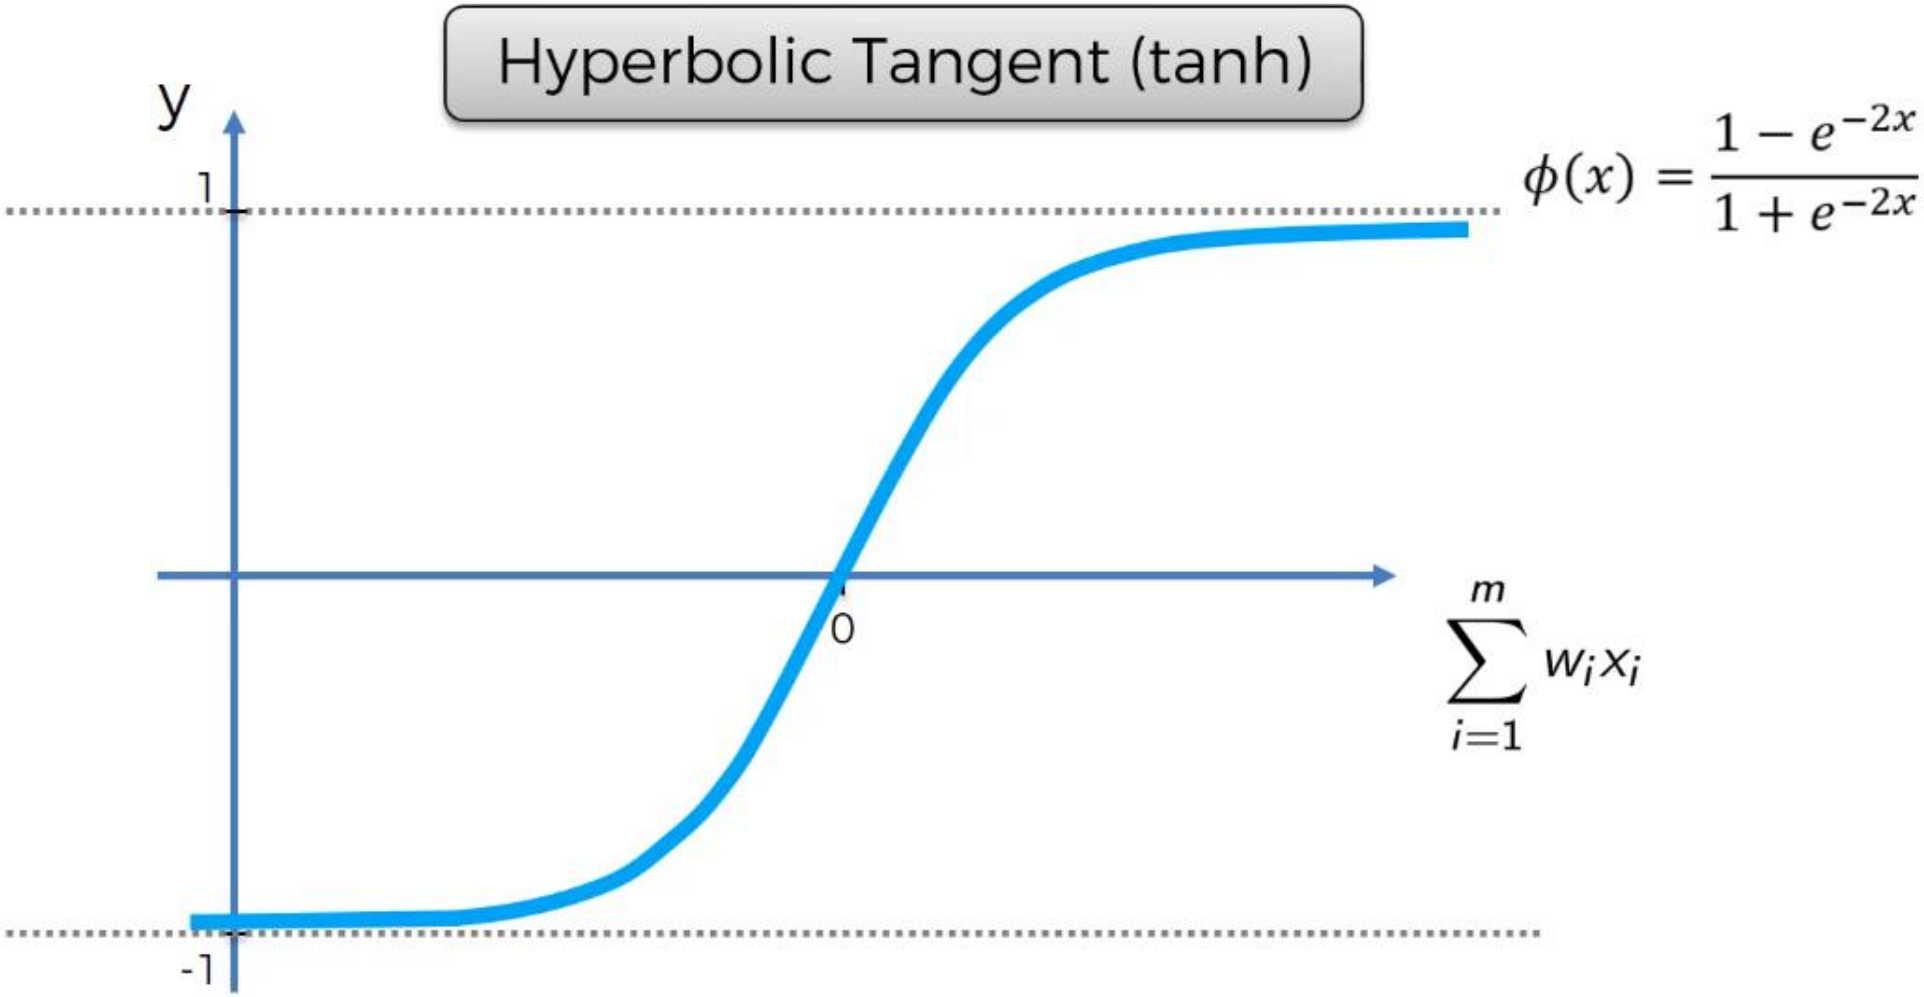

https://www.udemy.com/course/machinelearning/

The Activation Function

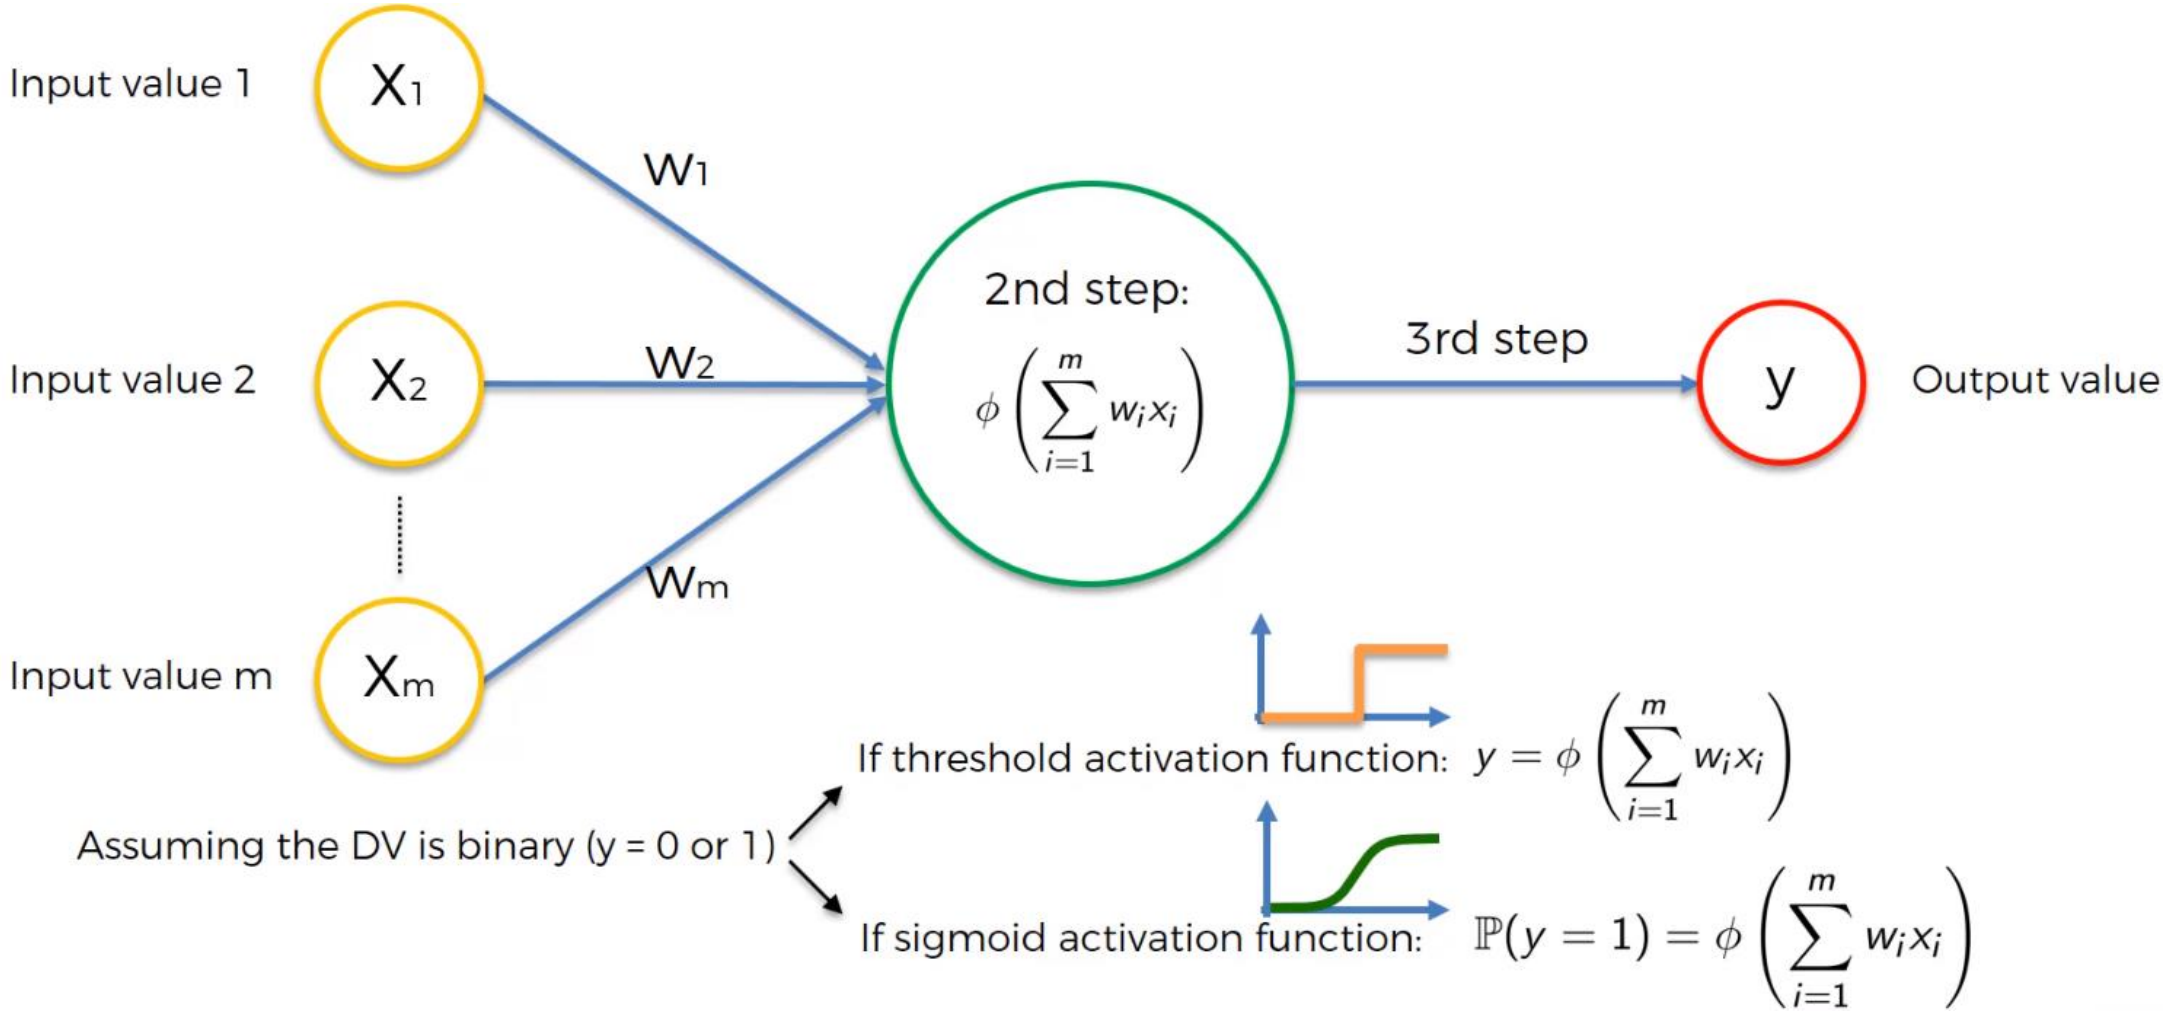

https://www.udemy.com/course/machinelearning/

The Activation Function

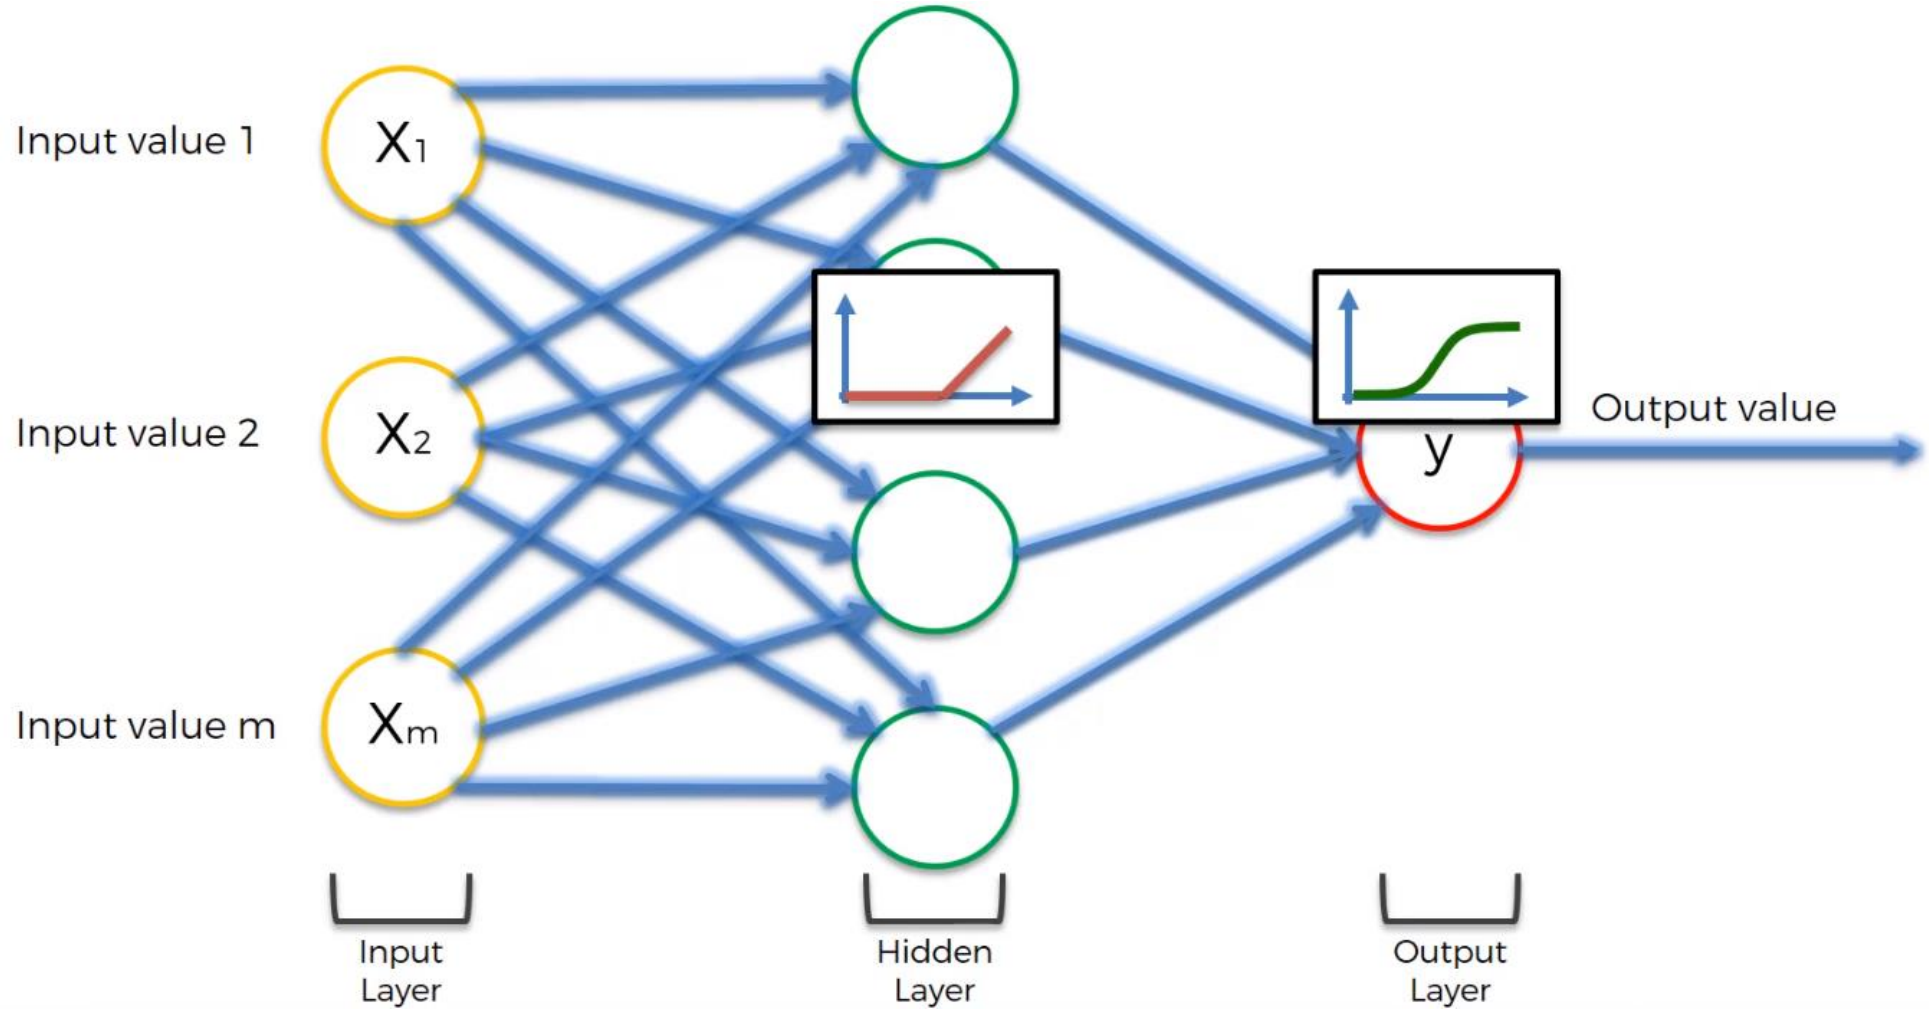

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

How do Neural Networks work?

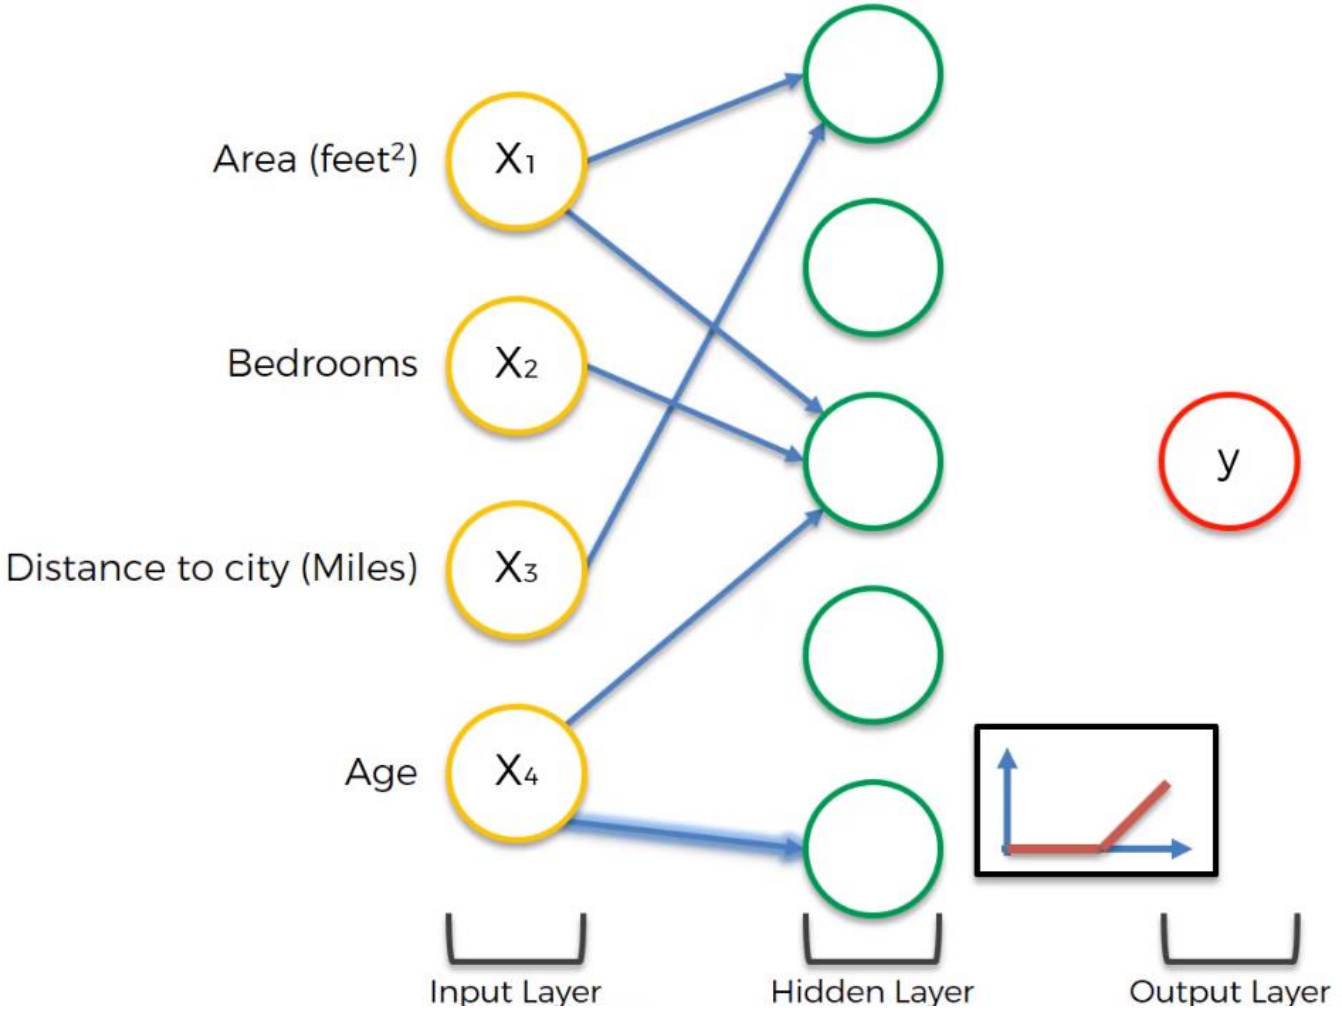

https://www.udemy.com/course/machinelearning/

How do Neural Networks work?

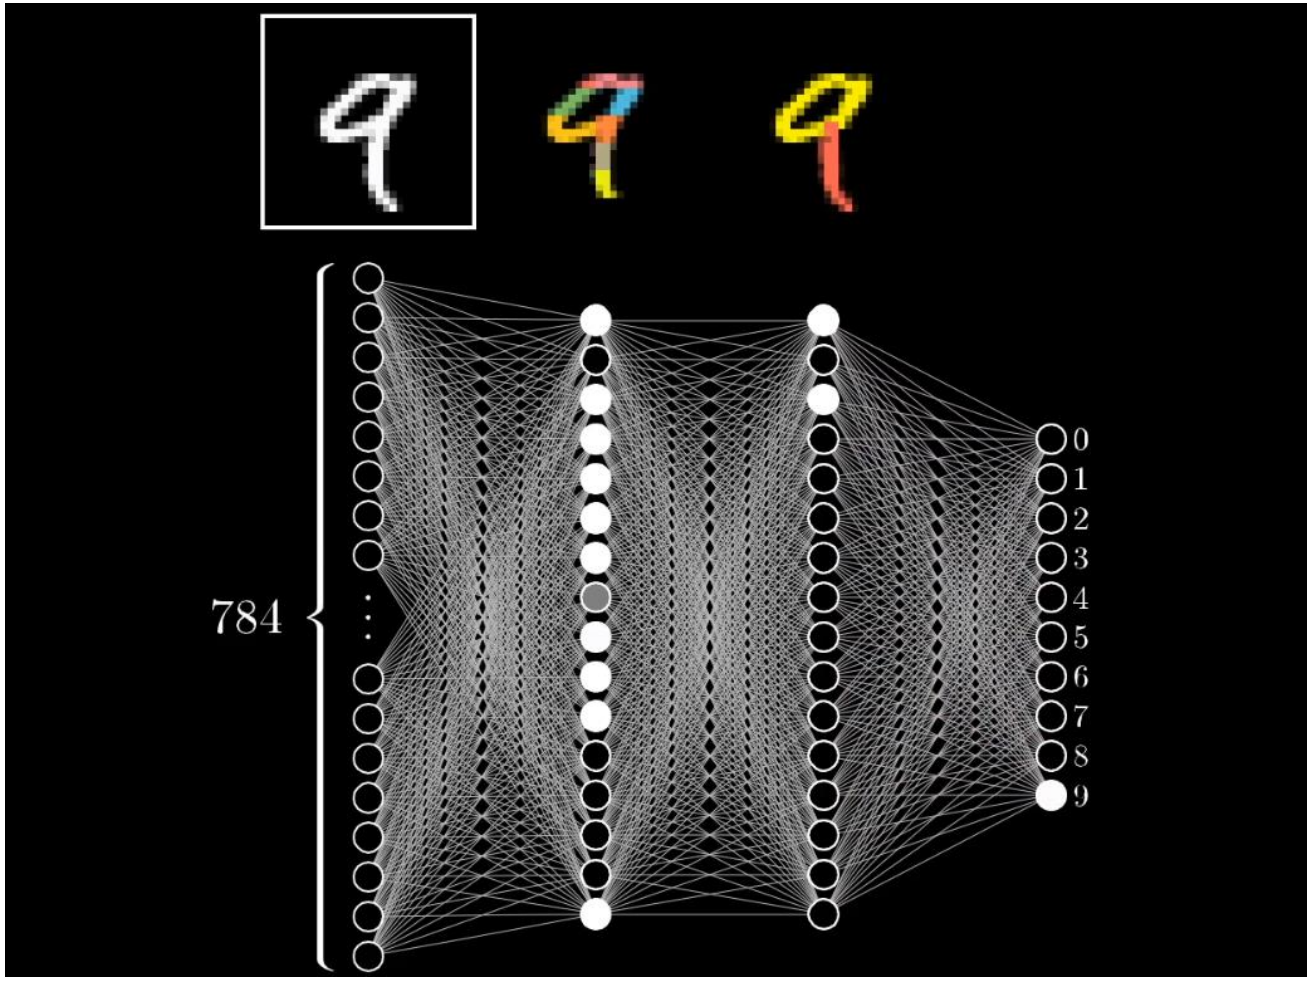

https://youtu.be/aircAruvnKk

How do Neural Networks work?

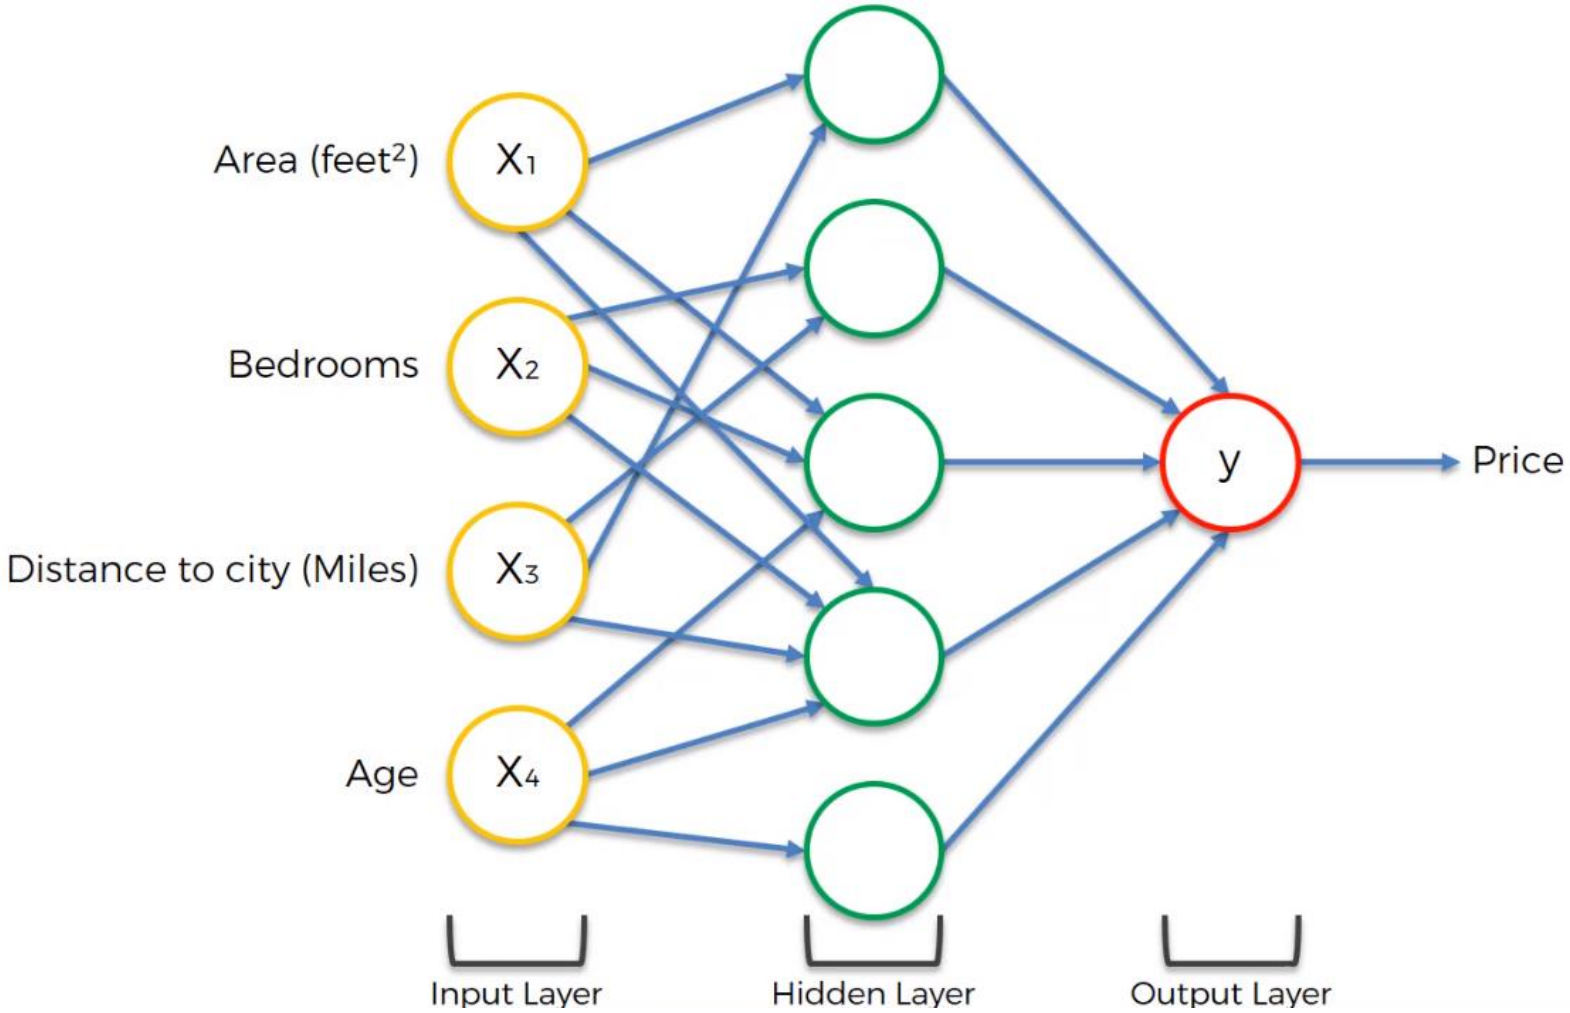

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

How do Neural Networks learn?

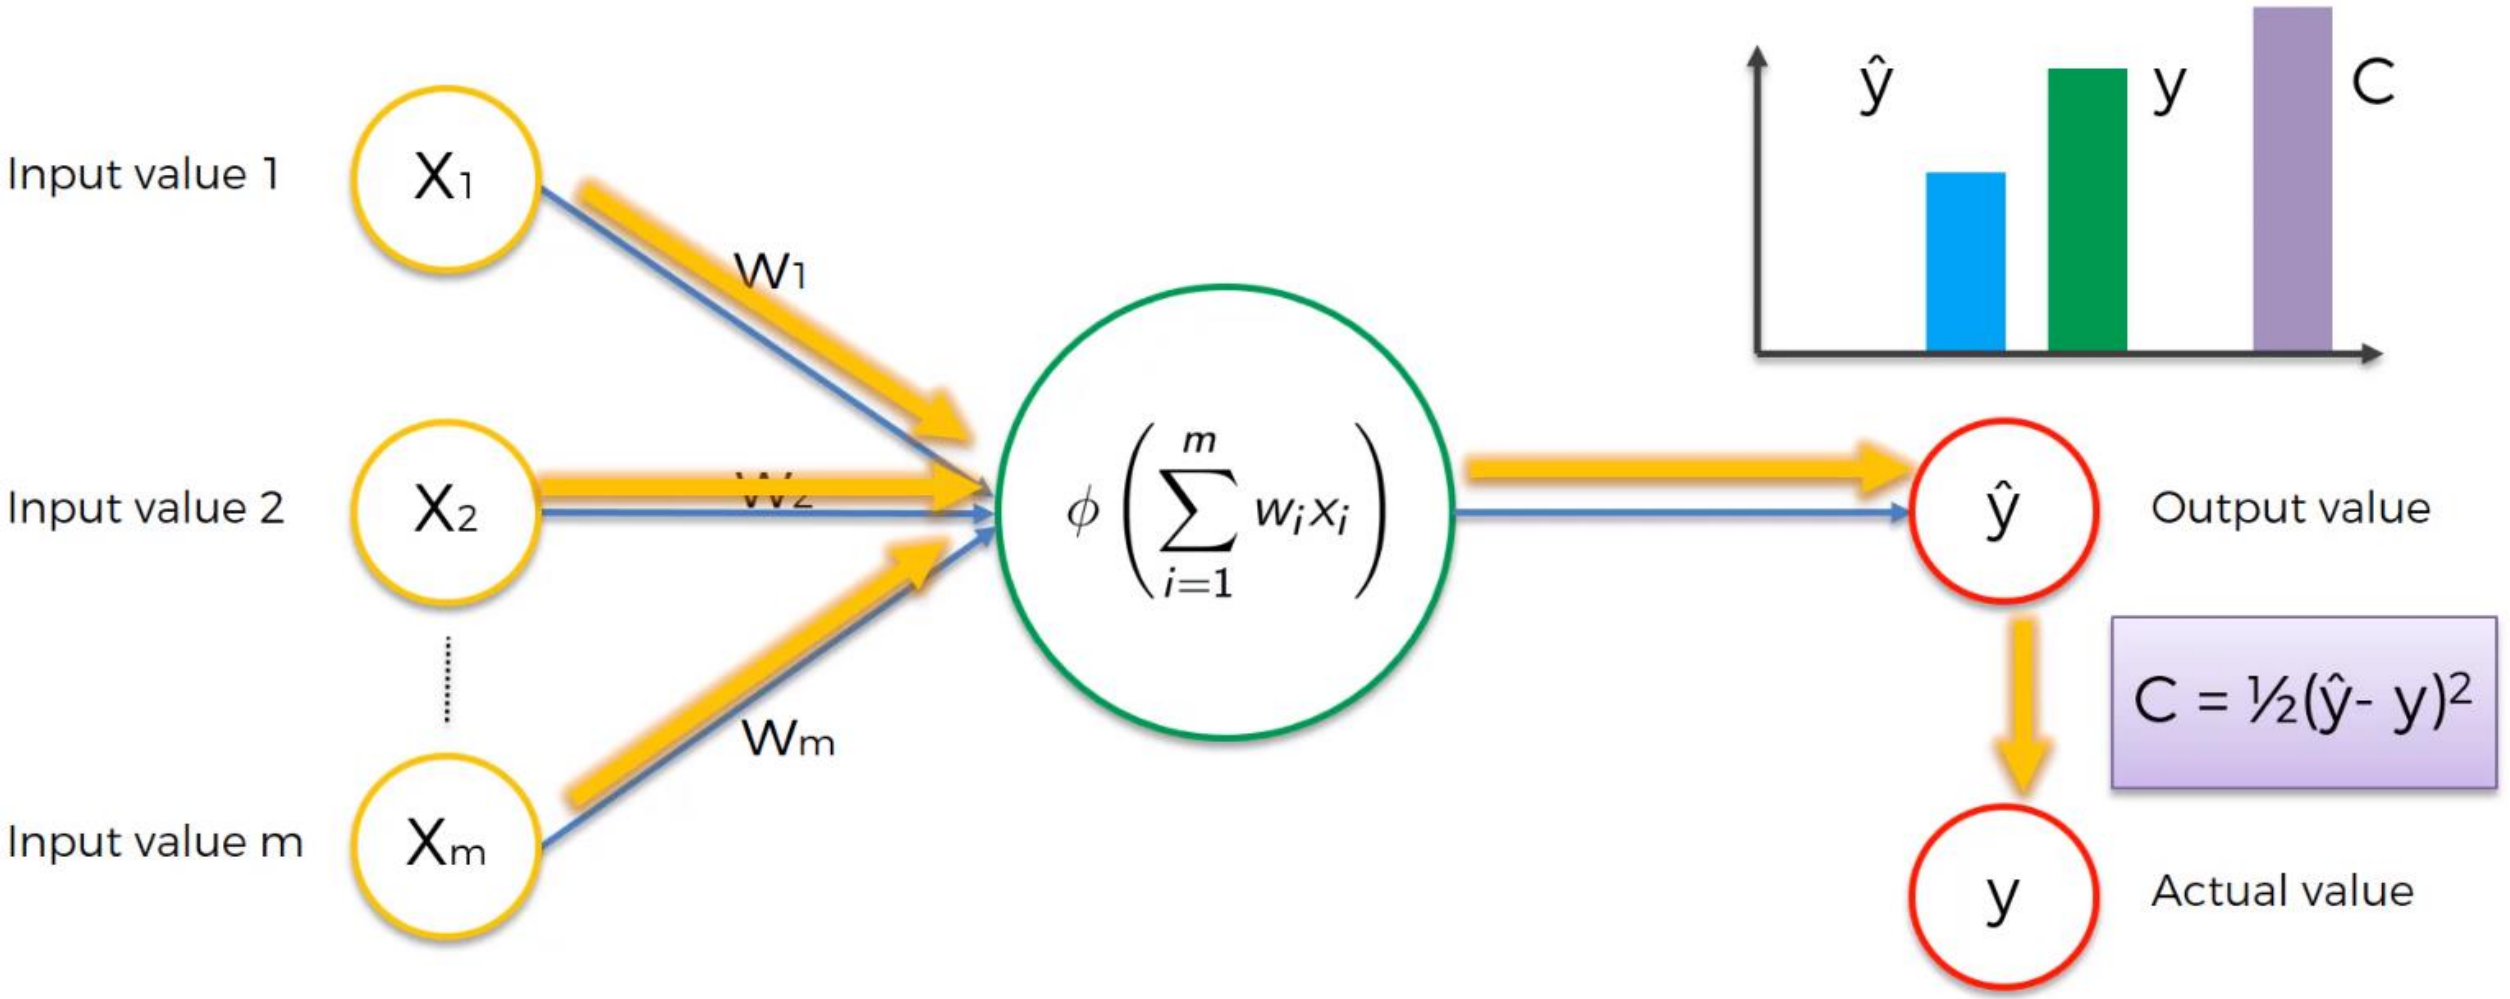

https://www.udemy.com/course/machinelearning/

How do Neural Networks learn?

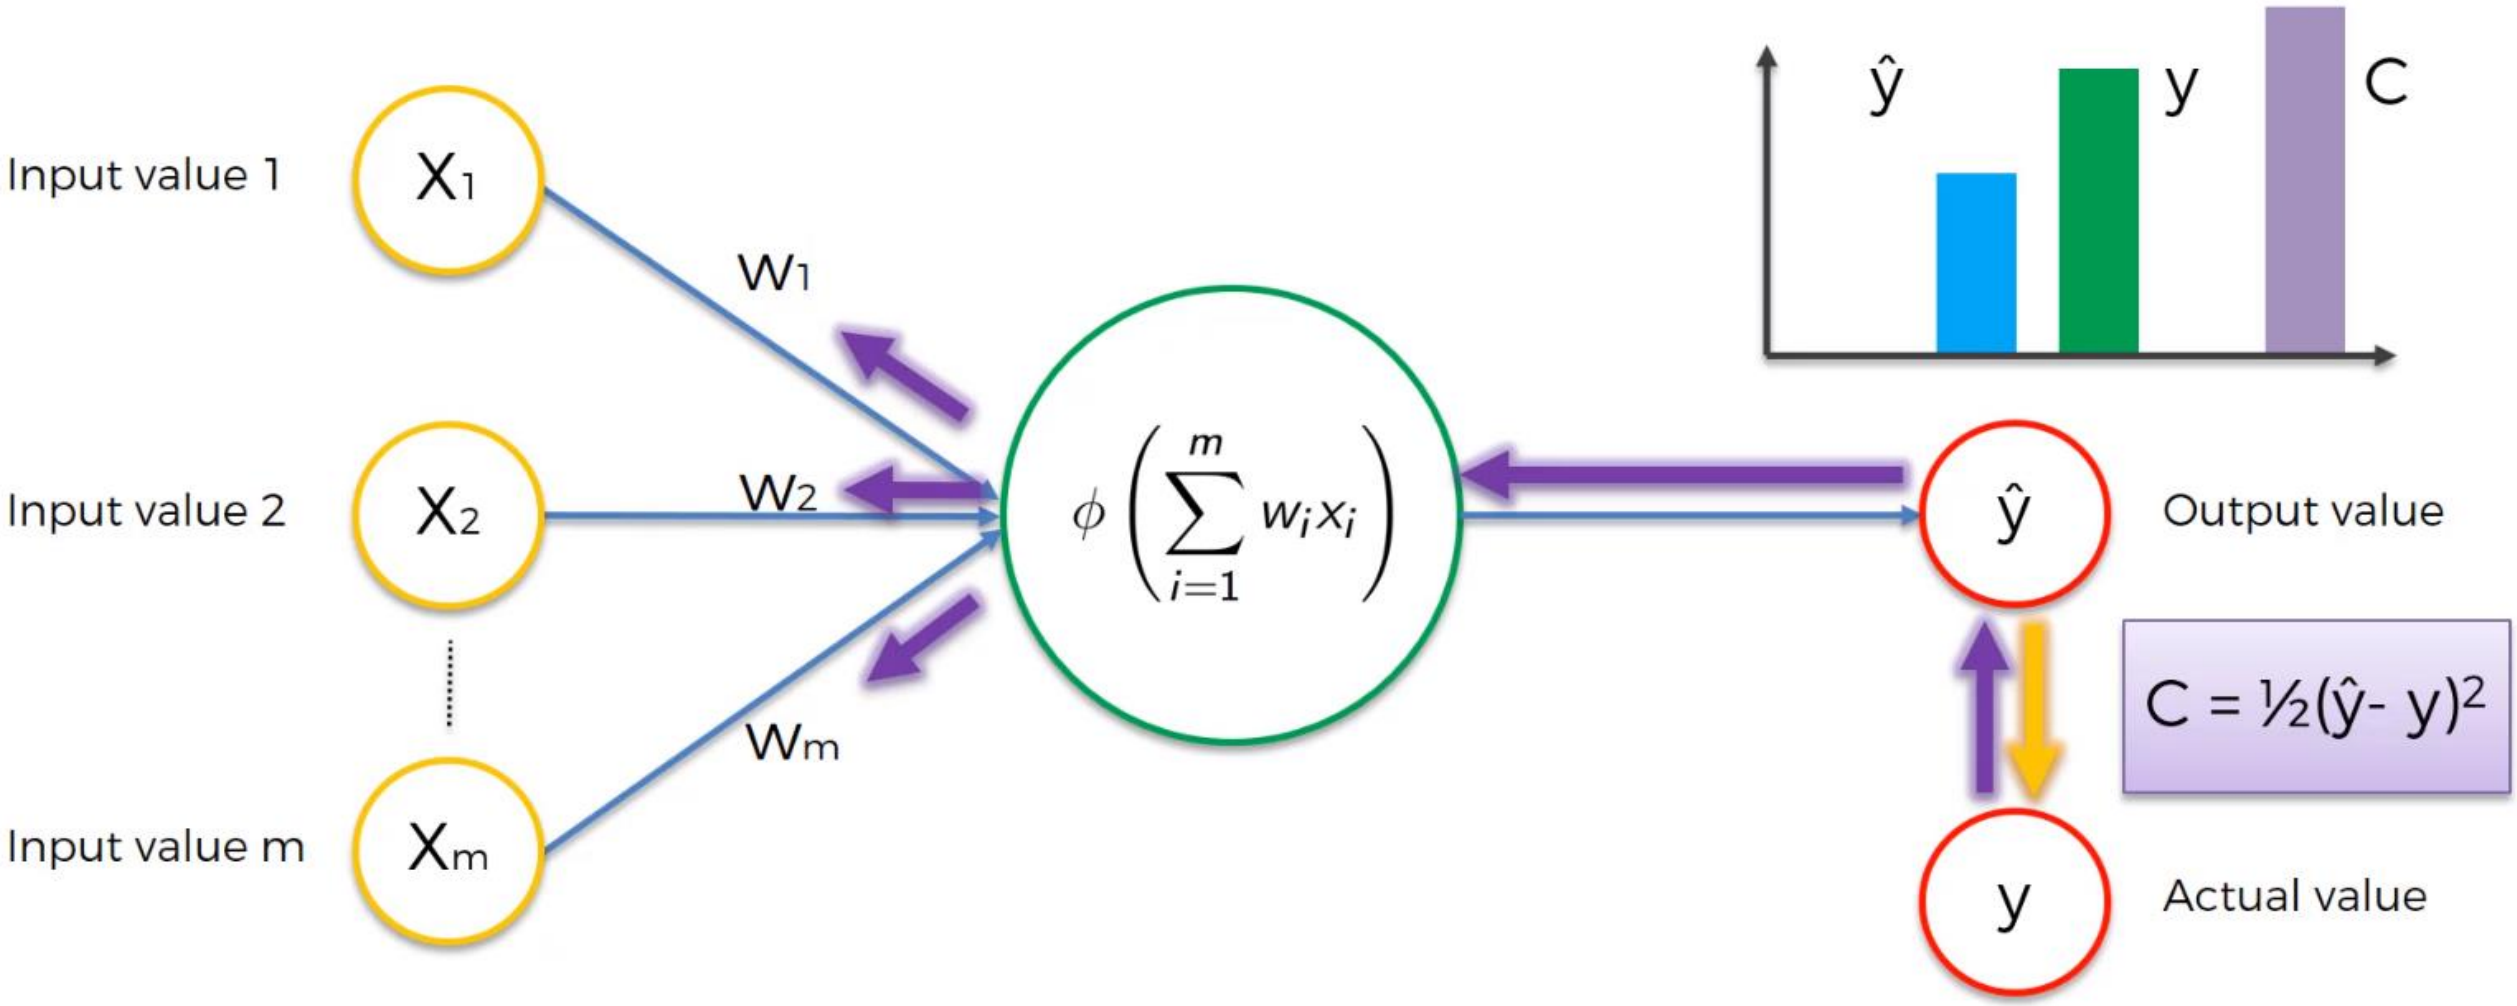

https://www.udemy.com/course/machinelearning/

How do Neural Networks learn?

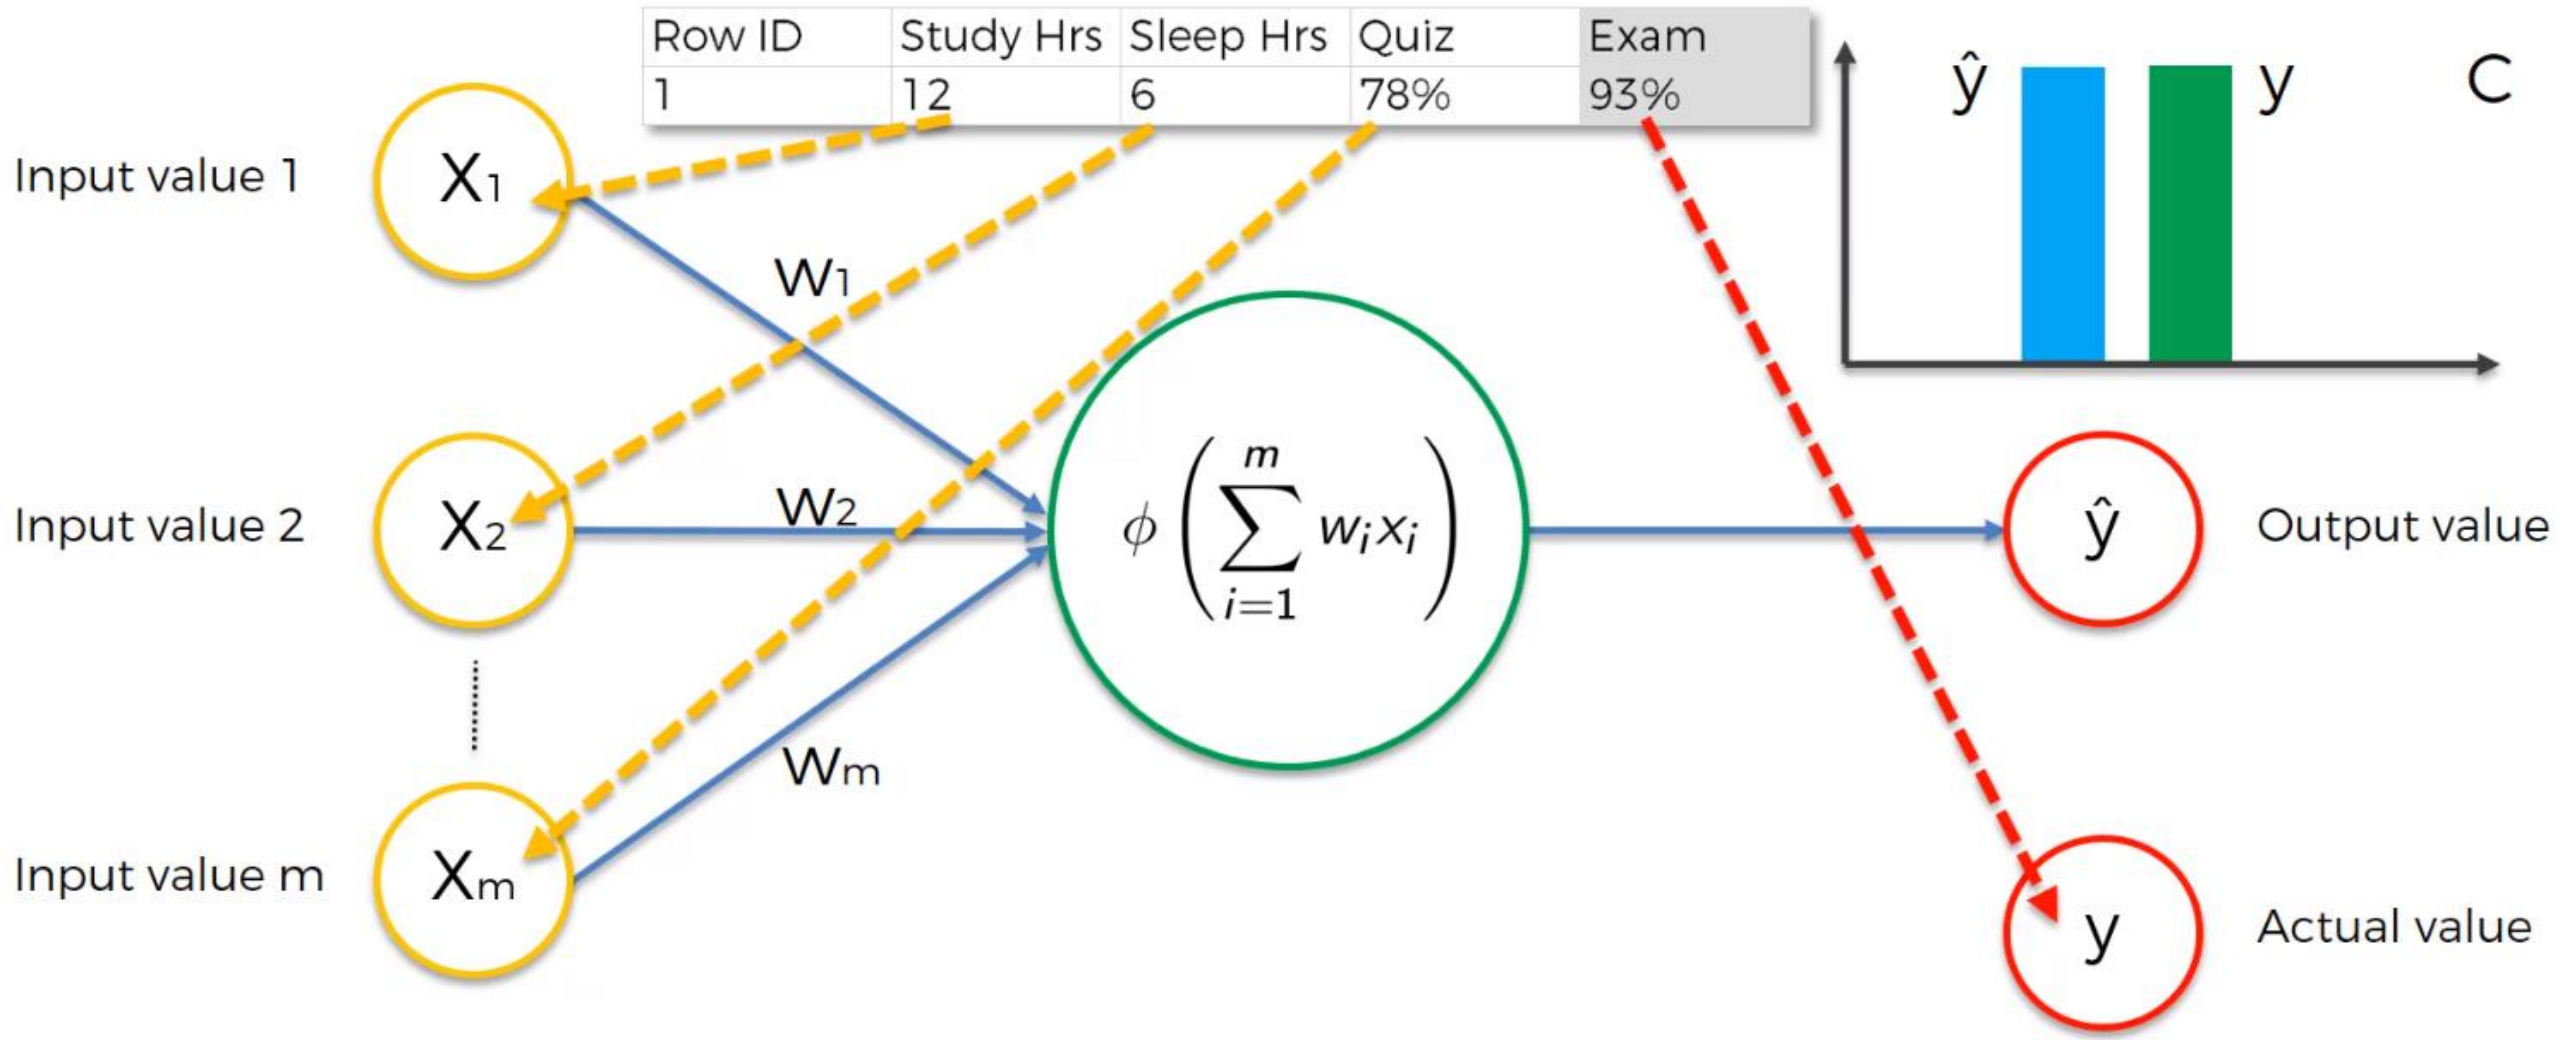

https://www.udemy.com/course/machinelearning/

How do Neural Networks learn?

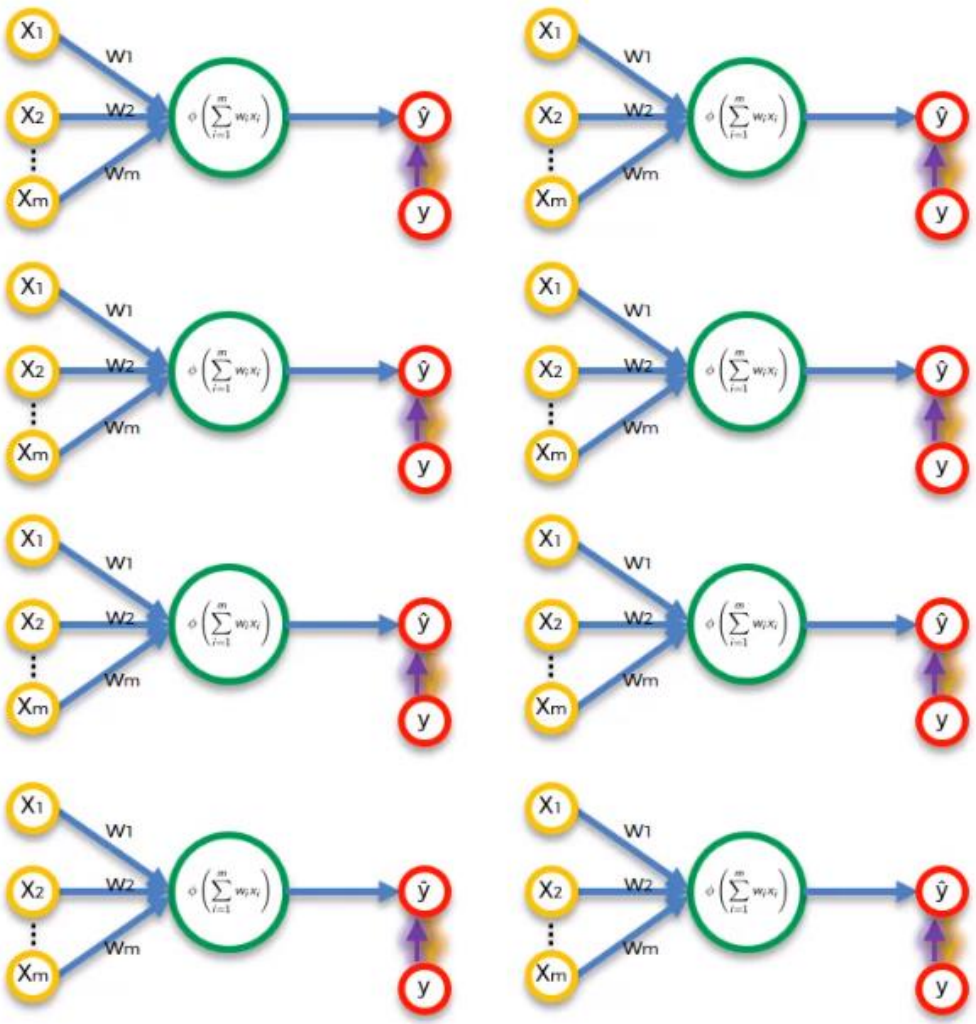

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

$$C = \sum \tfrac{1}{2}(\hat{y} - y)^2$$

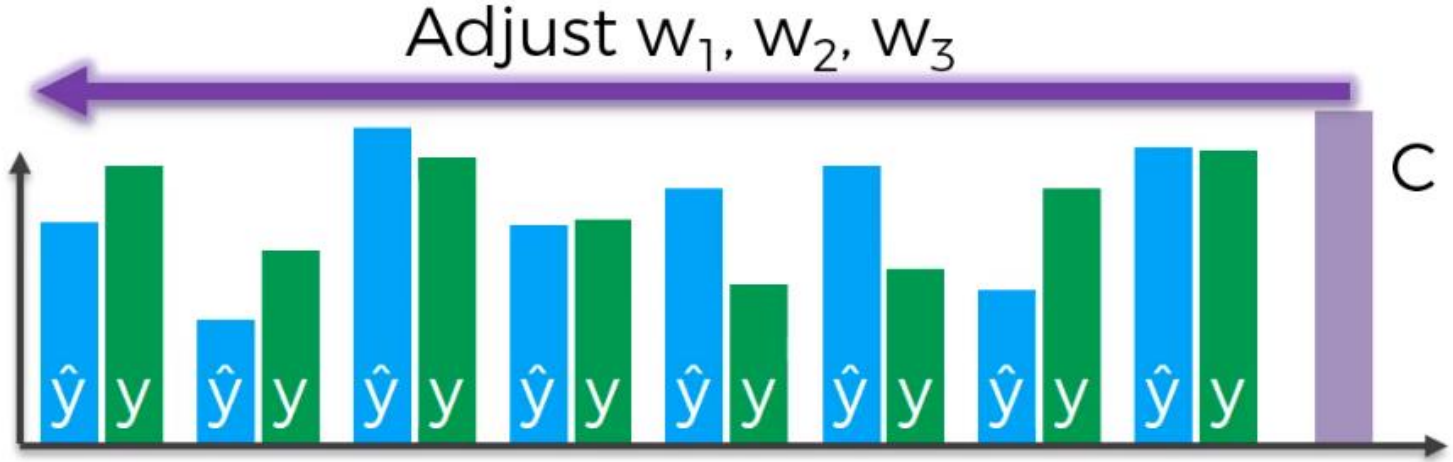

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

Gradient Descent

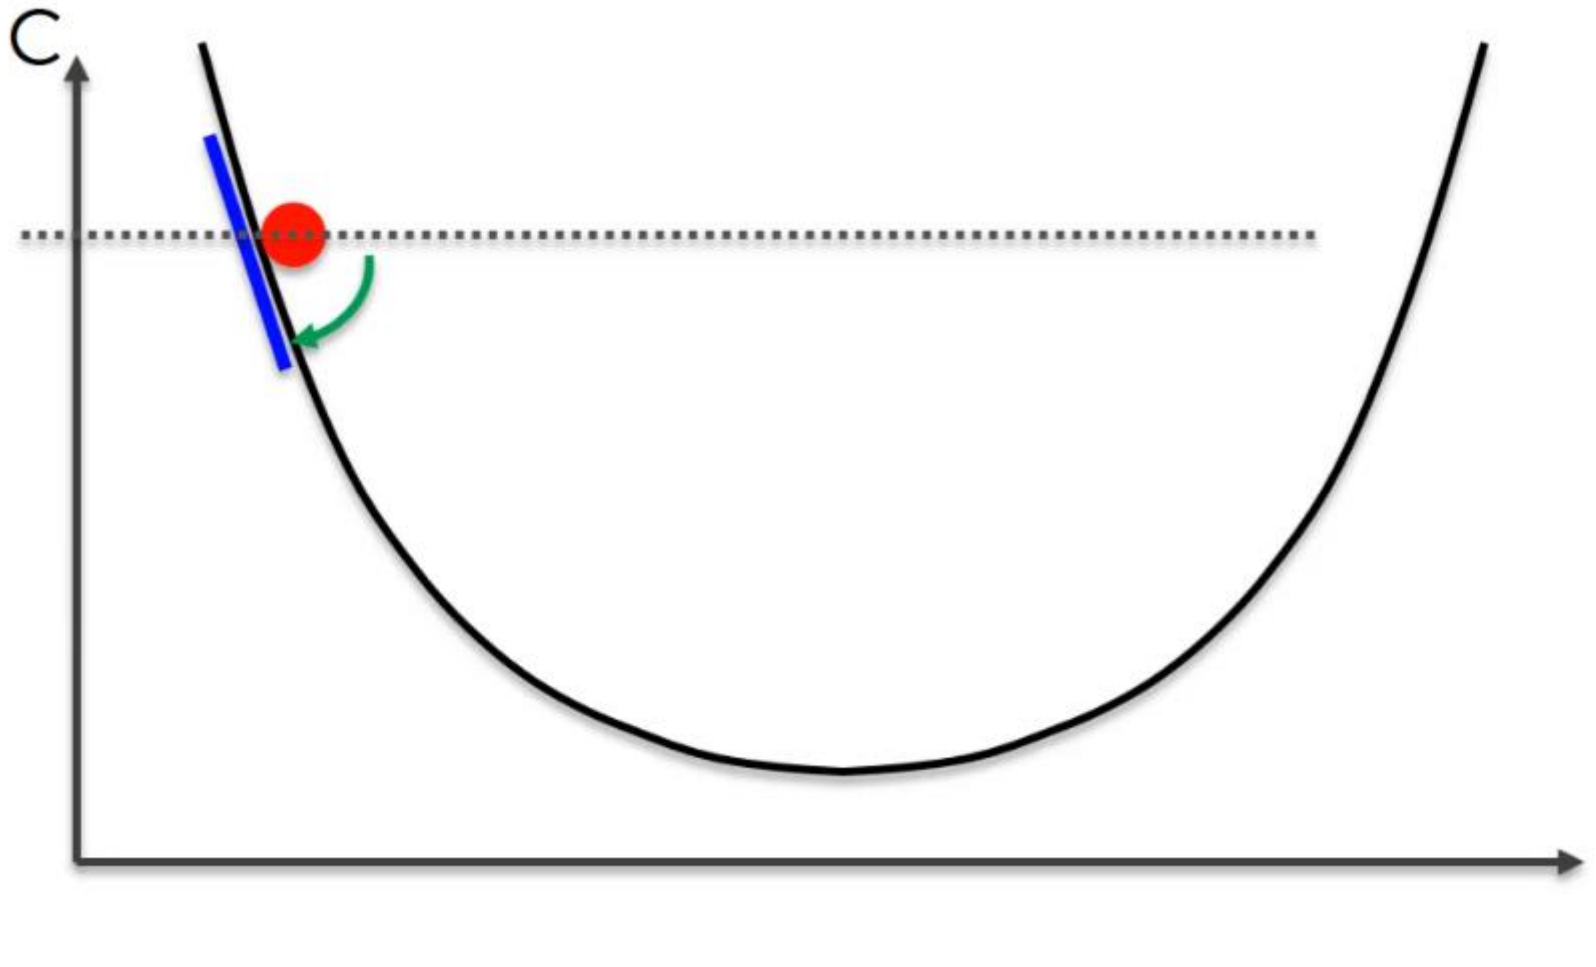

$$C = \frac{1}{2}(\hat{y} - y)^2$$

https://www.udemy.com/course/machinelearning/

Gradient Descent

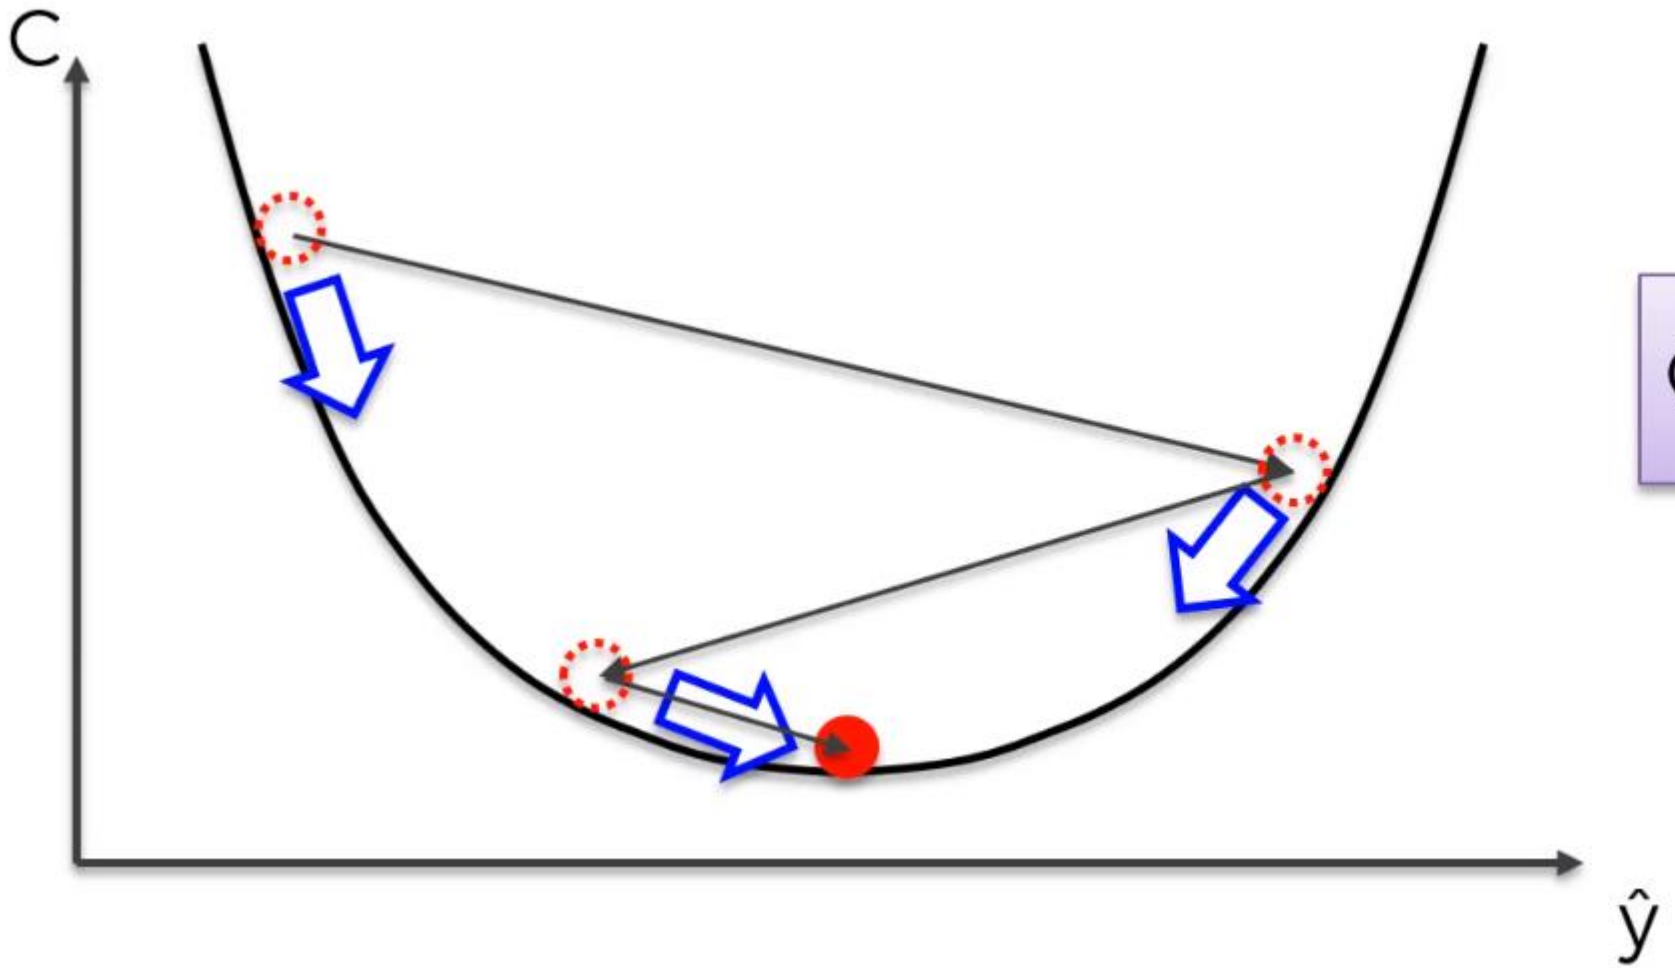

$$C = \frac{1}{2}(\hat{y} - y)^2$$

https://www.udemy.com/course/machinelearning/

Gradient Descent

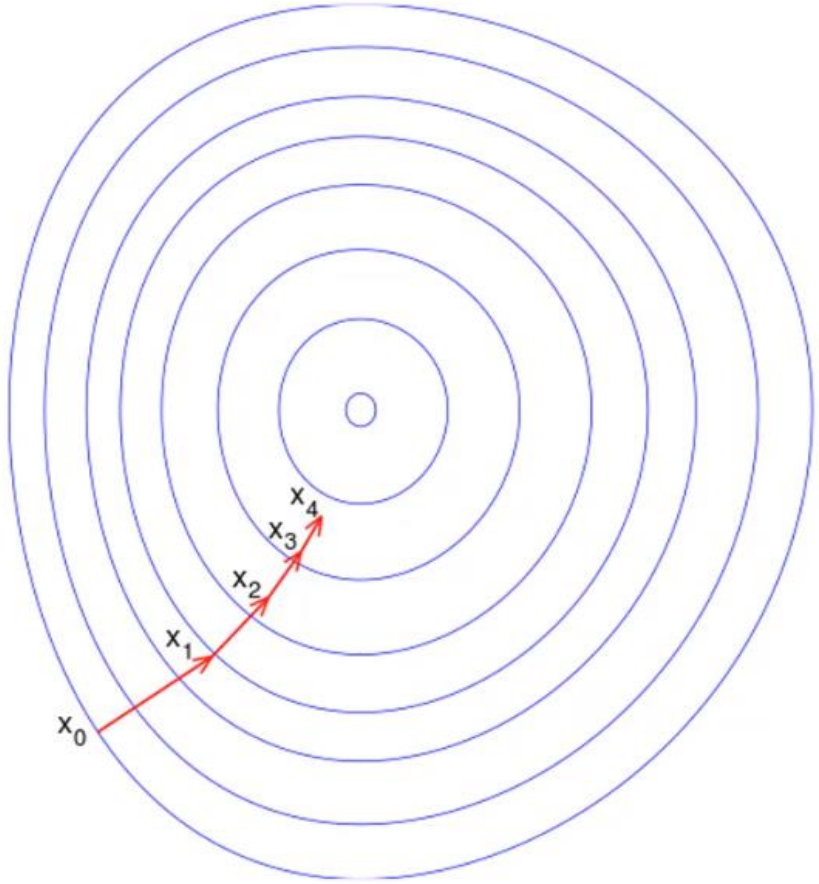

https://www.udemy.com/course/machinelearning/

Gradient Descent

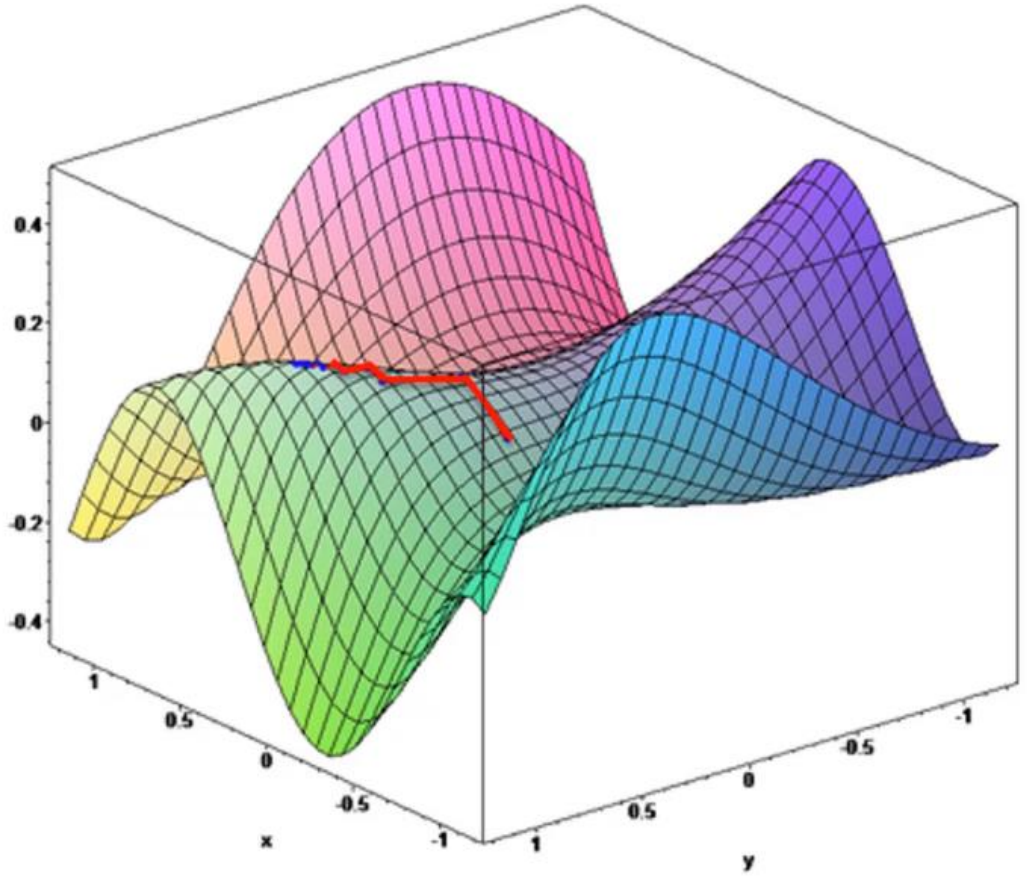

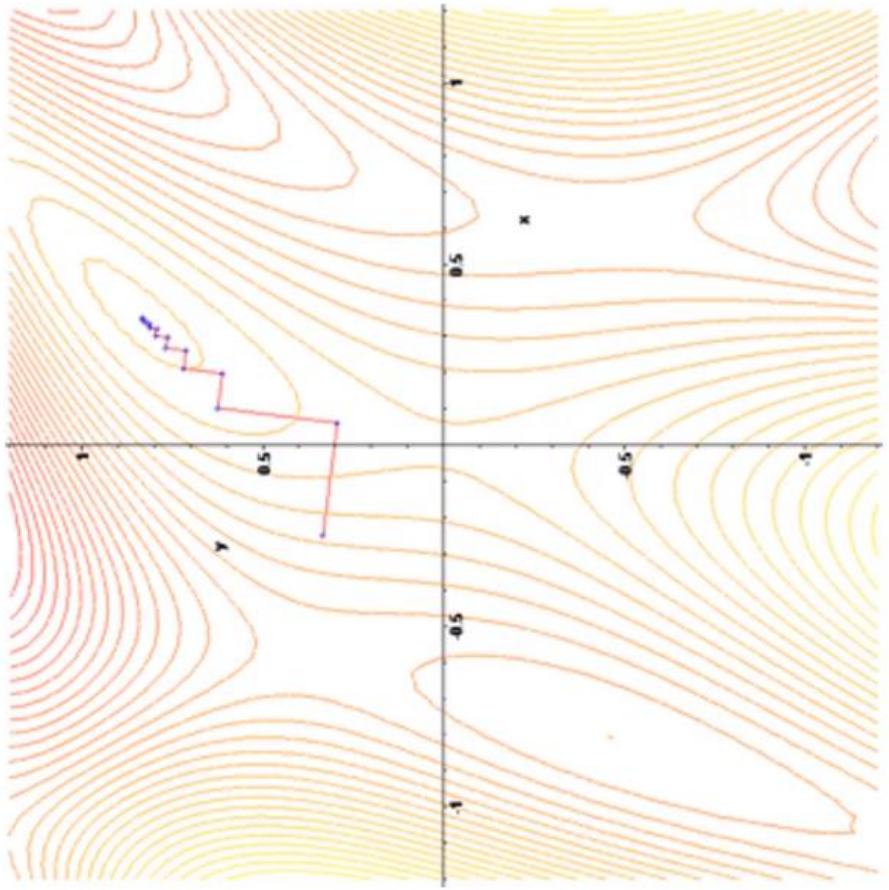

https://www.udemy.com/course/machinelearning/

Gradient Descent

Step 1: Take the derivative of the **Cost Function** for each parameter in it. In fancy Machine Learning lingo, take the **Gradient** of the **Loss Function**.

Step 2: Pick random values for the parameters.

Step 3: Plug the parameter values into the derivatives (**Gradient**).

Step 4: Calculate the Step Sizes: Step Size = **Slope * Learning Rate**

Step 5: Calculate the New Parameters:

New Parameter = Old Parameter – Step Size

Step 6: Go back to **Step 3** and repeat until **Step Size** is very small, or you reach the **Maximum Number of Steps**

https://youtu.be/sDv4f4s2SB8

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

Stochastic Gradient Descent

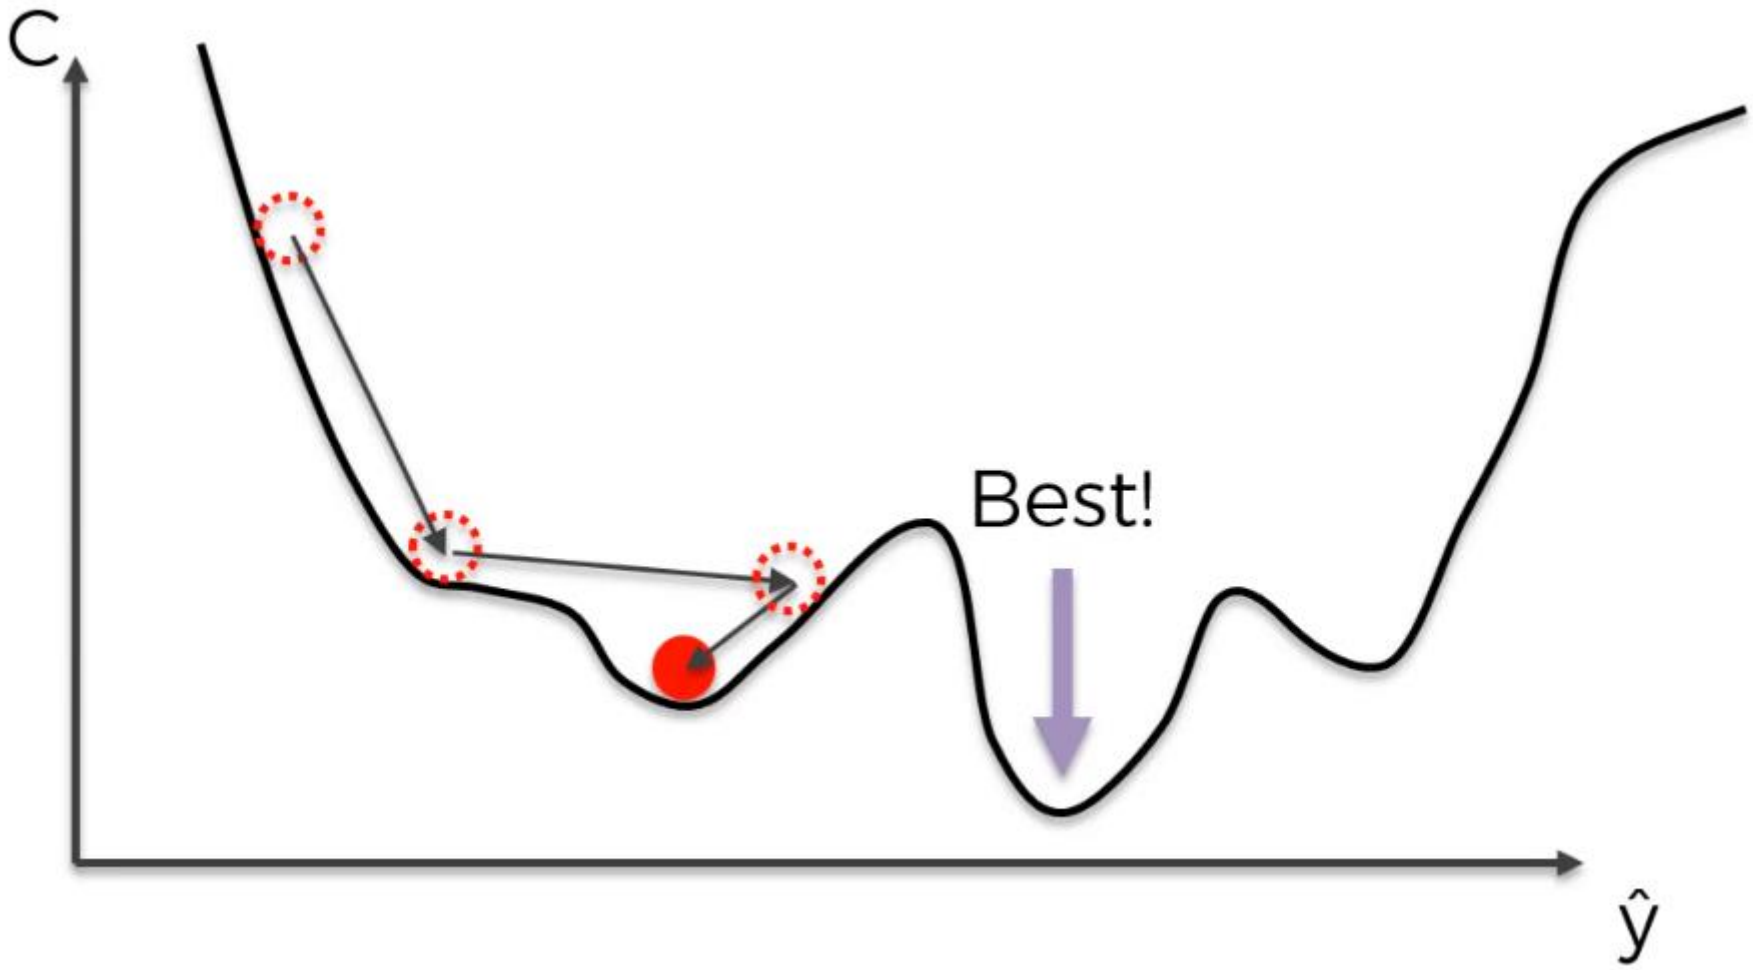

https://www.udemy.com/course/machinelearning/

Stochastic Gradient Descent

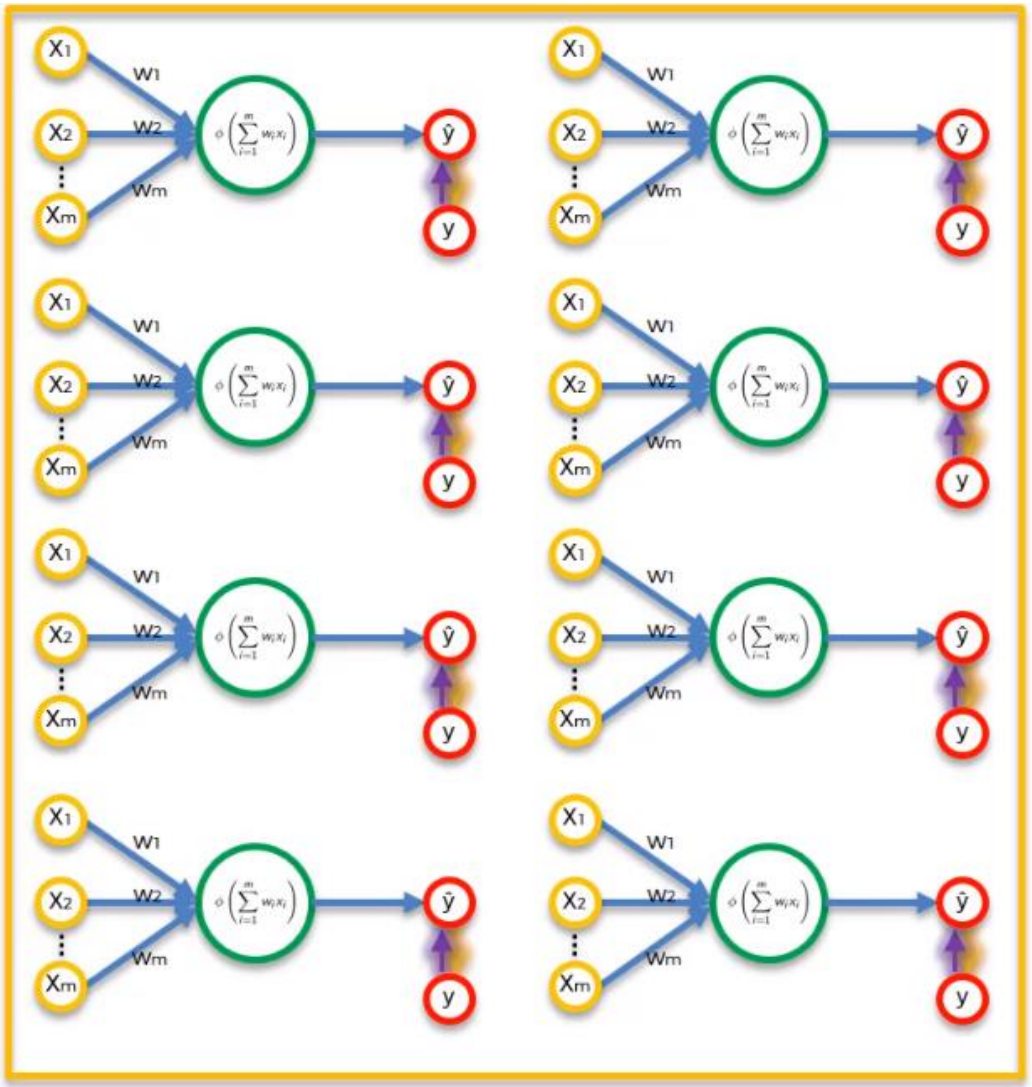

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

$$C = \sum \tfrac{1}{2}(\hat{y} - y)^2$$

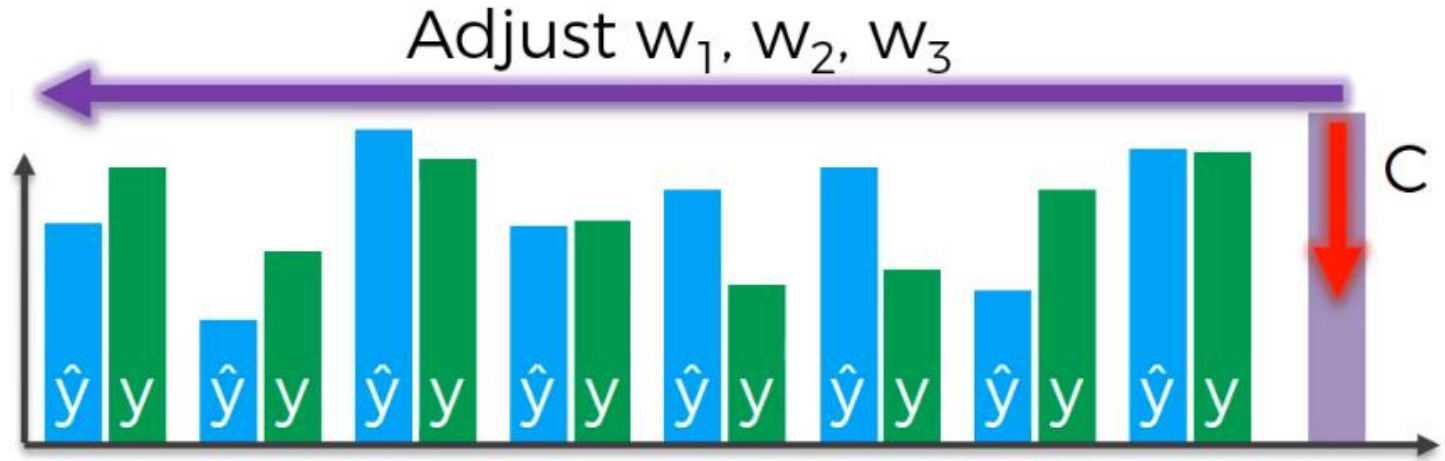

https://www.udemy.com/course/machinelearning/

Stochastic Gradient Descent

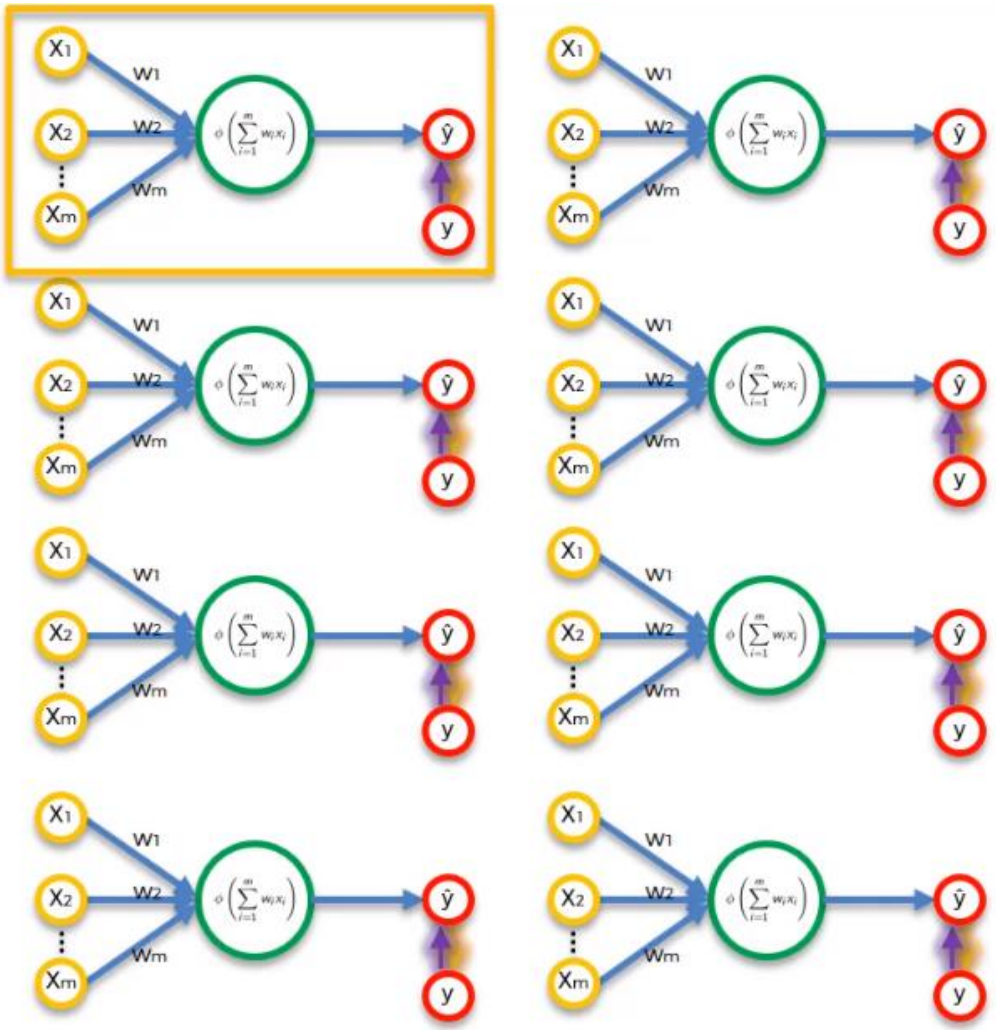

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

$$C = \sum \tfrac{1}{2}(\hat{y} - y)^2$$

https://www.udemy.com/course/machinelearning/

Stochastic Gradient Descent

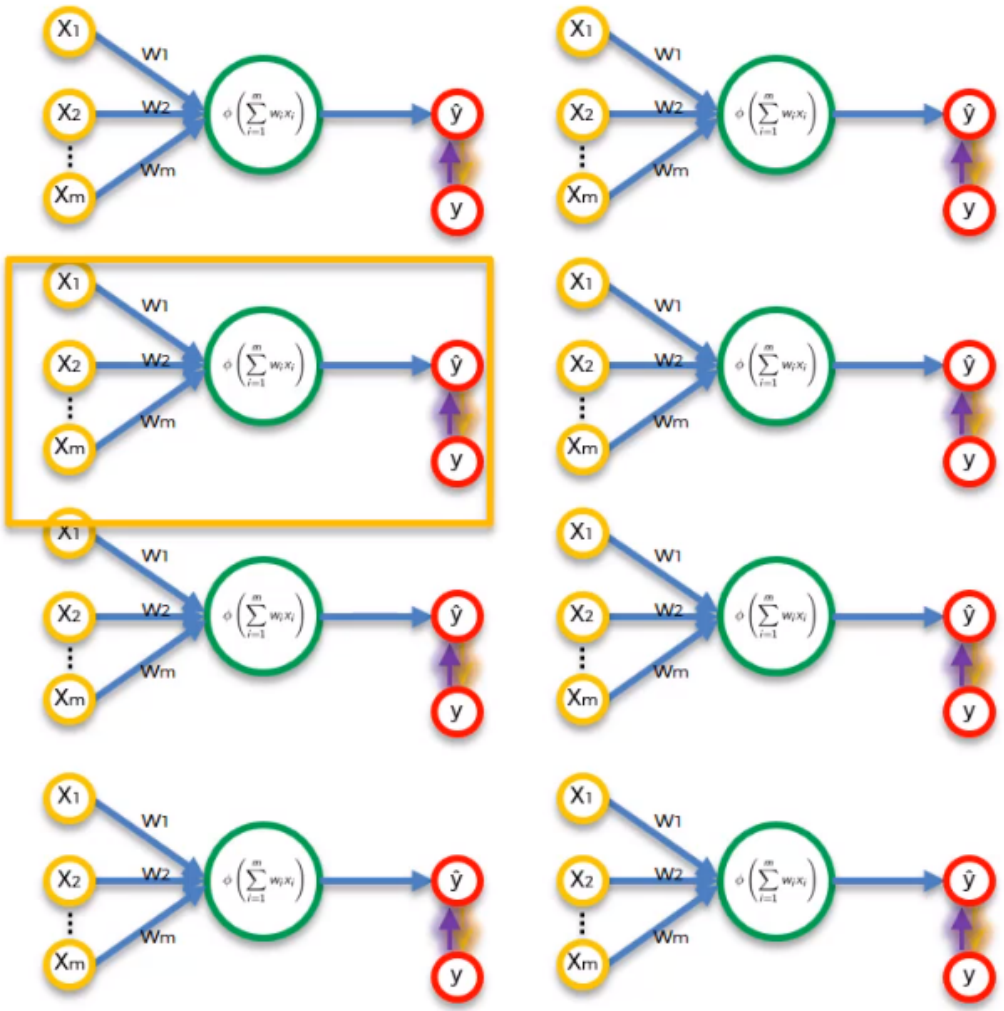

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

$$C = \sum \tfrac{1}{2}(\hat{y} - y)^2$$

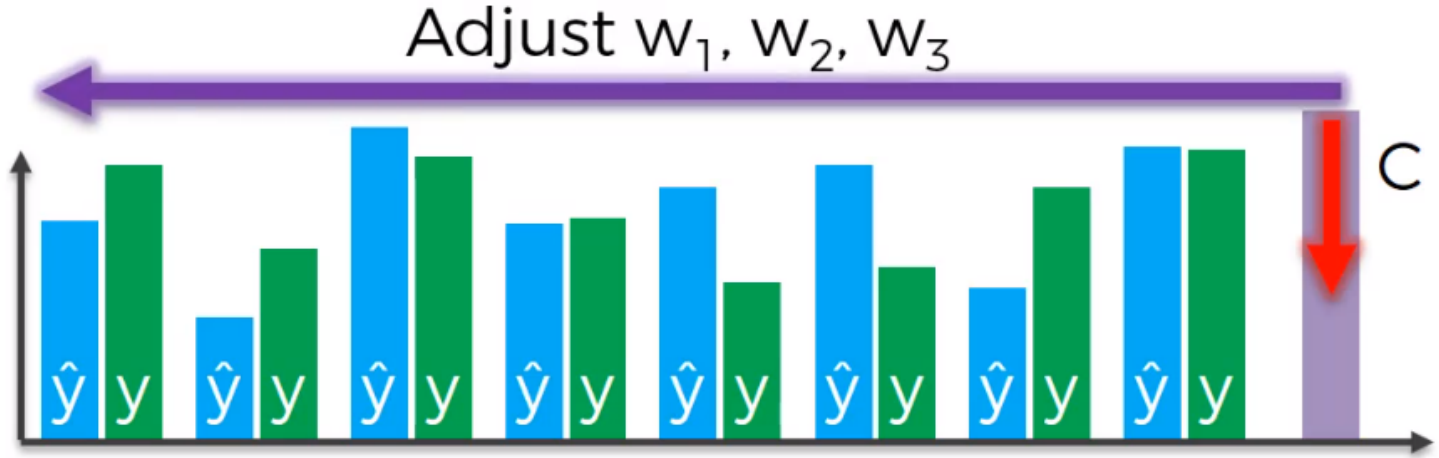

https://www.udemy.com/course/machinelearning/

Stochastic Gradient Descent

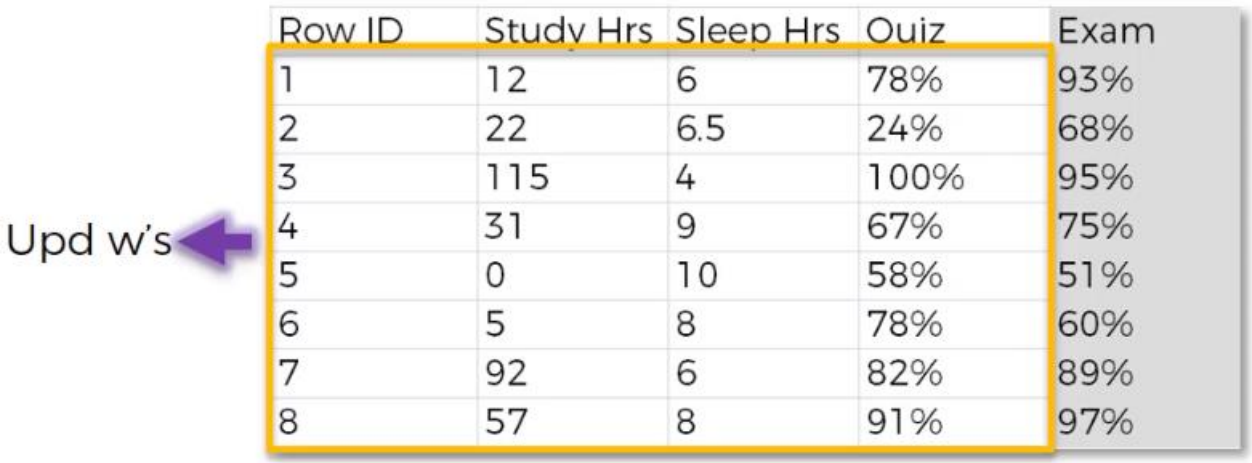

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Batch Gradient Descent

Stochastic Gradient Descent

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

Backpropagation

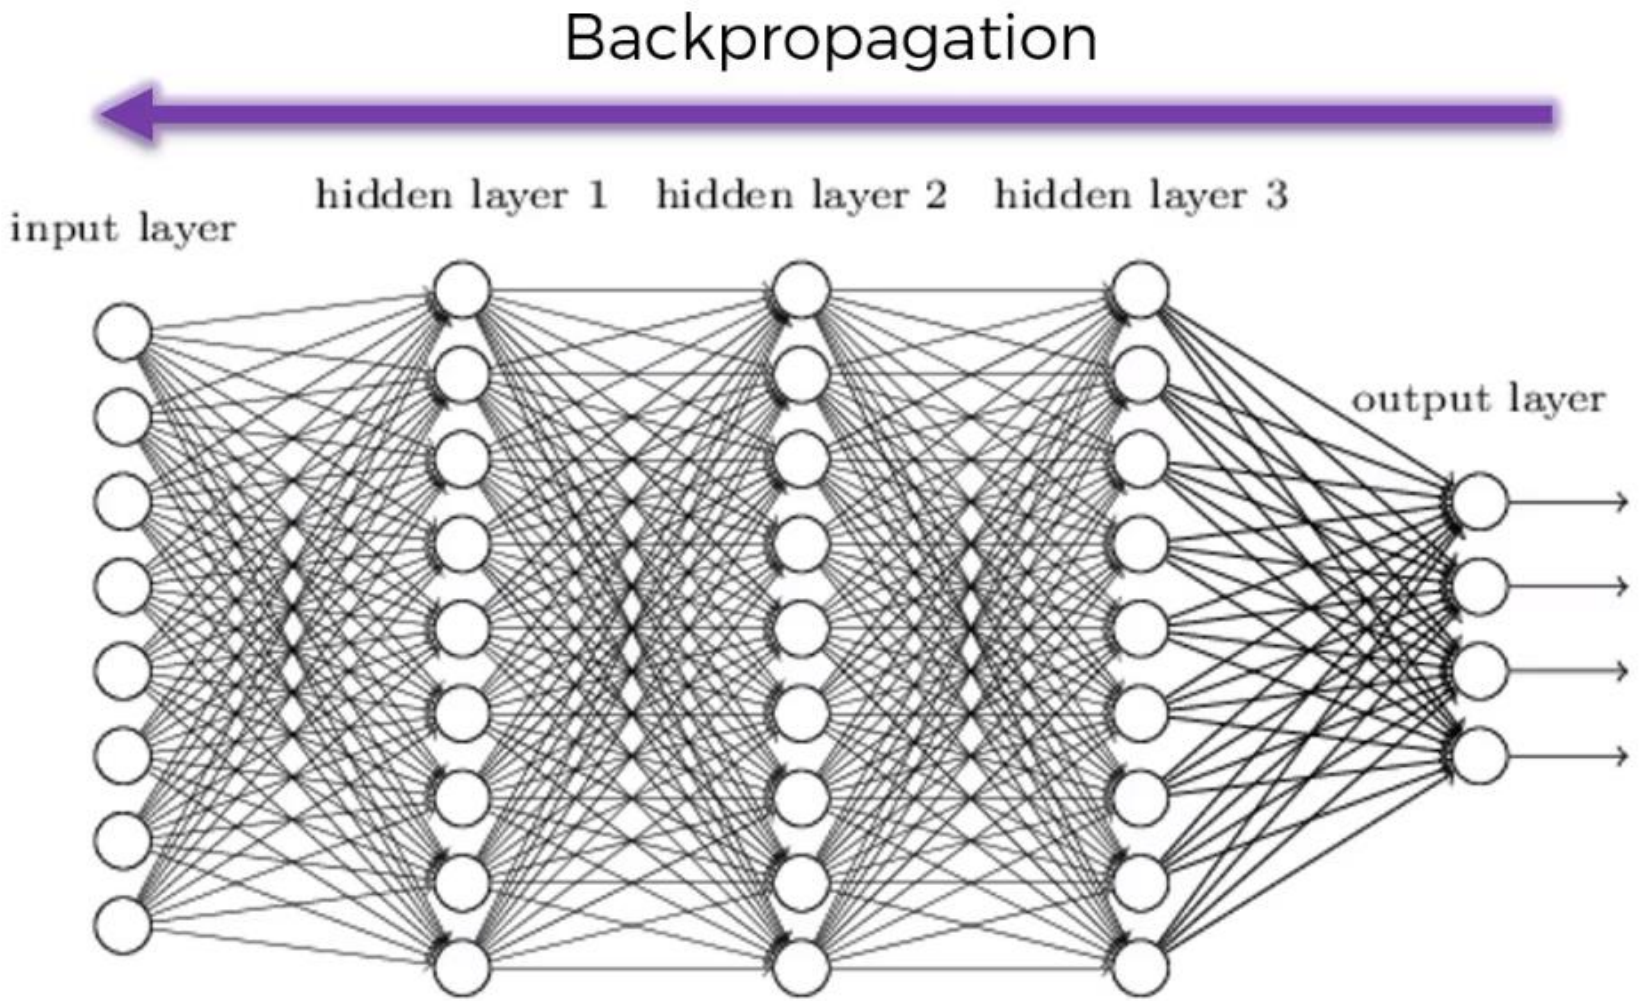

https://www.udemy.com/course/machinelearning/

Backpropagation

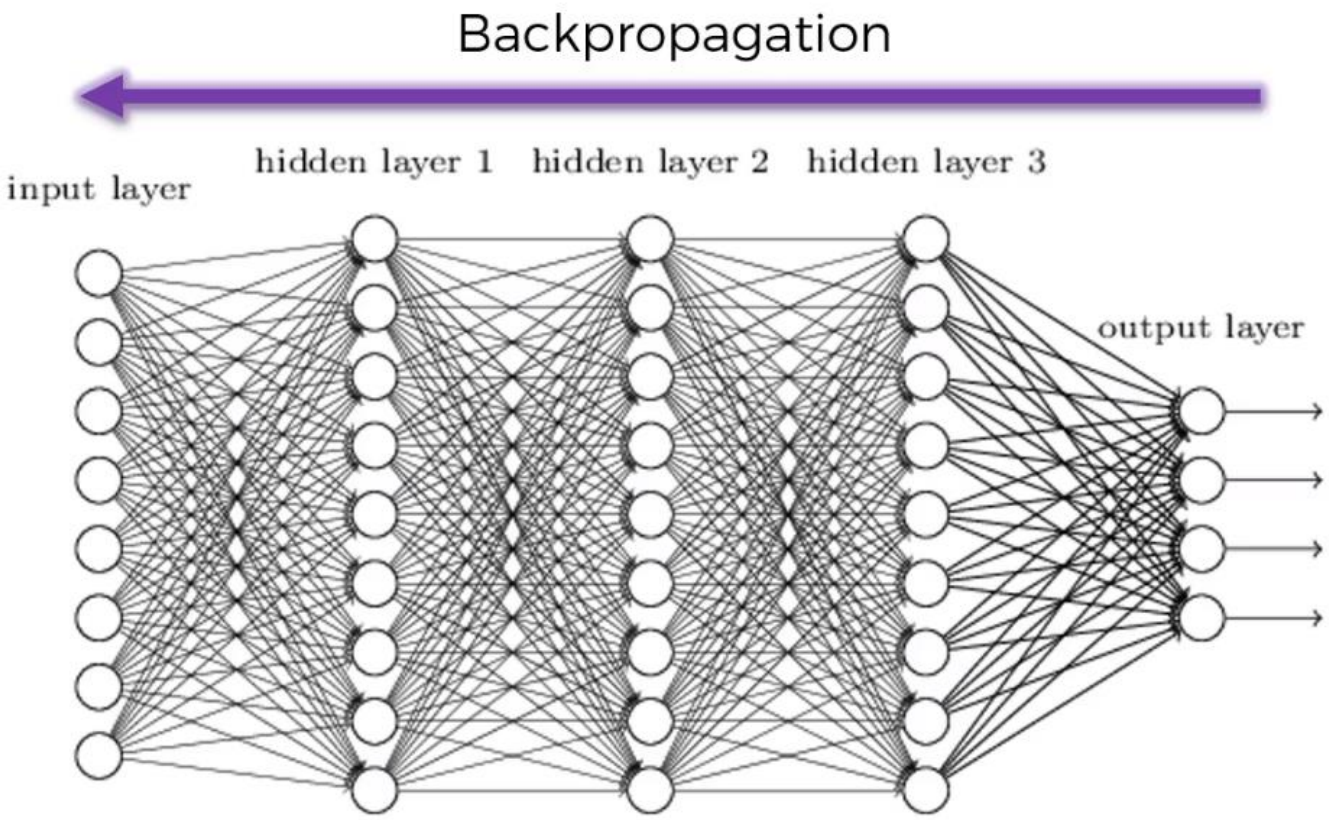

Gradient of the cost function

Input: 270 weights/biases
Output: 270 weights/biases nudges
Parameters: 1 number (the cost)

Cost function

Input: 270 weights/biases
Output: 1 number (the cost)
Parameters: Many, many, many training example

Neural Network function

Input: 8 numbers (features)
Output: 4 numbers
Parameters: 270 weights/biases

https://www.udemy.com/course/machinelearning/

Summary

- Overview
- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation
- Algorithm summary

https://www.udemy.com/course/machinelearning/

Algorithm summary

STEP 1: Randomly initialise the weights to small numbers close to 0 (but not 0).

STEP 2: Input the first observation of your dataset in the input layer, each feature in one input node.

STEP 3: Forward-Propagation: from left to right, the neurons are activated in a way that the impact of each neuron's activation is limited by the weights. Propagate the activations until getting the predicted result y.

STEP 4: Compare the predicted result to the actual result. Measure the generated error.

STEP 5: Back-Propagation: from right to left, the error is back-propagated. Update the weights according to how much they are responsible for the error. The learning rate decides by how much we update the weights.

STEP 6: Repeat Steps 1 to 5 and update the weights after each observation (Reinforcement Learning). Or: Repeat Steps 1 to 5 but update the weights only after a batch of observations (Batch Learning).

STEP 7: When the whole training set passed through the ANN, that makes an epoch. Redo more epochs.

https://www.udemy.com/course/machinelearning/

Algorithms used

- Neural Network architecture/topology: **Deep Feed Forward (DFF)**
- Activation function: **Rectifier** and **sigmoid**
- Training method/Optimizer:
  - **Backpropagation** to calculate the derivatives and
  - **Gradient Descent** which descends through the gradient, i.e. adjust the parameters
- Cost/loss function: $C = \sum \frac{1}{2}(\hat{y} - y)^2$